

# FinalProject

May 11, 2022

## 0.0.1 Note to reader

This notebook has been set up s.t. the activation functions and loss weight combination experiments can be performed.

## 0.0.2 Import modules + Check GPU

```
[23]: import torch
import torchvision
from torch import nn
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset
import os
from PIL import Image
from torchvision.io import read_image
from torchvision.transforms import Resize, Compose, ToTensor, Normalize
import numpy as np
import skimage
import matplotlib.pyplot as plt
import time
from skimage import io
from tqdm import tqdm
import scipy.ndimage
from torch.utils.tensorboard import SummaryWriter
from pathlib import Path

# for SSIM
import math

# for beta selection
import random

[4]: print("GPU available: {}".format(torch.cuda.is_available()))
print("Device: {}".format(torch.cuda.get_device_name(0)))
```

```
GPU available: True
Device: Tesla K80
```

### 0.0.3 Data Generation

```
[5]: def isotropic_diffusion(img, niter=1, kappa=50, gamma=0.1, voxelspacing=None):

    # initialize output array
    out = np.array(img, dtype=np.float32, copy=True)

    # set default voxel spacing if not supplied
    if voxelspacing is None:
        voxelspacing = tuple([1.] * img.ndim)

    # initialize some internal variables
    deltas = [np.zeros_like(out) for _ in range(out.ndim)]

    time = 0

    results_pixels = []
    results_dIdt = []
    results_time = []

    results_pixels.append(out.astype(img.dtype))
    results_time.append(time)
    #results_dIdt.append(np.zeros_like(out))

    for iter in tqdm(range(niter)):
        # calculate the diffs
        for i in range(out.ndim):
            slicer = [slice(None, -1) if j == i else slice(None) for j in
↪range(out.ndim)]
            diff_local = np.diff(out, axis=i)
            deltas[i][tuple(slicer)] = diff_local

        matrices = [delta for delta, spacing in zip(deltas, voxelspacing)]

        # second derivative
        for i in range(out.ndim):
            slicer = [slice(1, None) if j == i else slice(None) for j in
↪range(out.ndim)]
            matrices[i][tuple(slicer)] = np.diff(matrices[i], axis=i)

        dIdt = np.sum(matrices, axis=0)
        #print(dIdt)

        # update the image
        out += gamma * (dIdt)
```

```

        time += gamma

        results_dIdt.append(dIdt.astype(img.dtype))
        if iter < niter - 1:
            results_pixels.append(out.astype(img.dtype))
            results_time.append(time)

    return results_pixels, results_dIdt, results_time

def get_mgrid(sidelen=256, dim=2):

    '''Generates a flattened grid of (x,y,...) coordinates in a range of -1 to 1.
    ↪1.
    sidelen: int
    dim: int'''

    tensors = tuple(dim * [torch.linspace(-1, 1, steps=sidelen)])
    mgrid = torch.stack(torch.meshgrid(*tensors), dim=-1)
    mgrid = mgrid.reshape(-1, dim)

    return mgrid

class ImageFitting(Dataset):

    def __init__(self, img_path, niter):

        self.transform = Compose([
            Resize(256),
            ToTensor(),
            Normalize(torch.Tensor([0.5]), torch.Tensor([0.5]))
        ])
        self.coords = get_mgrid()

        print("-----Generating Data-----")
        self.base_img = io.imread(img_path)
        self.imgs_pixels, self.imgs_dIdt, self.imgs_time = ↪
        isotropic_diffusion(self.base_img, niter=niter, kappa=50, gamma=1/(niter+1))

        print("-----Finished-----")

        self.len = len(self.imgs_pixels)

    def __len__(self):

        return self.len

    def __getitem__(self, idx):

```

```

image = self.imgs_pixels[idx]
image = self.transform(Image.fromarray(image))

pixels = image.permute(1, 2, 0).view(-1, 1)
step_val = torch.full((self.coords.size(0),1), self.imgs_time[idx])

model_input = torch.cat((self.coords, step_val), 1)

# Compute gradient and laplacian
grads_x = scipy.ndimage.sobel(image.numpy(), axis=1).squeeze(0)[...,  
↪None]
grads_y = scipy.ndimage.sobel(image.numpy(), axis=2).squeeze(0)[...,  
↪None]
grads_x, grads_y = torch.from_numpy(grads_x), torch.from_numpy(grads_y)

grads = torch.stack((grads_x, grads_y), dim=-1).view(-1, 2)
laplace = scipy.ndimage.laplace(image.numpy()).squeeze(0)[..., None]
laplace = torch.from_numpy(laplace).view(-1, 1)

dIdt = torch.from_numpy(self.imgs_dIdt[idx])
dIdt = dIdt.permute(0,1).view(-1)

return model_input, {'pixels':pixels, 'grads':grads, 'laplace':laplace,  
↪'dIdt':dIdt}

```

#### 0.0.4 Loss Calculation

```

[6]: def computeJacobianFull(x, outputs, create_graph):

    dy_dx = torch.autograd.grad(outputs=outputs, inputs=x, grad_outputs=torch.  
↪ones_like(outputs),
                                retain_graph=True, create_graph=create_graph, allow_unused=True)[0]

    dy_dx = dy_dx.view(outputs.size(0), outputs.size(1), dy_dx.size(2))

    return dy_dx

def computeLaplaceFull(x, jacobian, create_graph):

    div = 0
    for j in range(jacobian.size(-1)):

        dy_dx2 = torch.autograd.grad(outputs=jacobian[:, :, j], inputs=x,  
↪grad_outputs=torch.ones_like(jacobian[:, :, j]),
                                    retain_graph=True, create_graph=create_graph)[0][..., j:j+1]

```

```

        div += dy_dx2

    return div

def calcLoss(coords, model_output, gt):

    pixel_loss = ((model_output - gt['pixels'])**2).mean()

    gradients = computeJacobianFull(coords, model_output, create_graph=True)
    grad_loss = ((gradients[:, :, :-1] - gt['grads']).pow(2).sum(-1)).mean()

    laplacian = computeLaplaceFull(coords, gradients[:, :, :-1],
    ↪create_graph=False)
    laplacian_loss = ((laplacian - gt['laplace'])**2).mean()

    dIdt_loss = ((gradients[:, :, -1] - gt['dIdt'])**2).mean()

    pixel_ssim = mean_ssim(gt['pixels'][0].cpu().view(1, 256, 256).detach(),
    ↪model_output[0].cpu().view(1, 256, 256).detach(), val_range=255)

    grad_ssim = mean_ssim(gt['grads'][0].norm(dim=-1).cpu().view(1, 256, 256).
    ↪detach(), gradients[0][:, :, -1].norm(dim=-1).cpu().view(1, 256, 256).detach(),
    ↪val_range=255)

    laplacian_ssim = mean_ssim(gt['laplace'][0].cpu().view(1, 256, 256).
    ↪detach(), laplacian[0].cpu().view(1, 256, 256).detach(), val_range=255)

    dIdt_ssim = mean_ssim(gt['dIdt'][0].cpu().view(1, 256, 256).float().
    ↪detach(), gradients[0][:, -1].cpu().view(1, 256, 256).detach(), val_range=255)

    return pixel_loss, grad_loss, laplacian_loss, dIdt_loss, pixel_ssim,
    ↪grad_ssim, laplacian_ssim, dIdt_ssim

```

### 0.0.5 SSIM (Structural Similarity Index Measure)

original SSIM paper: <https://www.cns.nyu.edu/pub/eero/wang03-reprint.pdf>

code source: <https://github.com/pranjaldata/SSIM-PyTorch>

explanation: <https://medium.com/srm-mic/all-about-structural-similarity-index-ssim-theory-code-in-pytorch-6551b455541e>

```

[7]: def gaussian(window_size=11, sigma=1.5):
    """
    Generates a list of Tensor values drawn from a gaussian distribution with
    ↪standard
    deviation = sigma and sum of all elements = 1.

    Length of list = window_size

```

```

"""
    gauss = torch.Tensor([math.exp(-(x - window_size//2)**2/float(2*sigma**2))])
    ↪for x in range(window_size)]
    ↪return gauss/gauss.sum()

```

```

[8]: def create_window(window_size=11, channel=1):

    # Generate an 1D tensor containing values sampled from a gaussian
    ↪distribution
    _1d_window = gaussian(window_size=window_size, sigma=1.5).unsqueeze(1)

    # Converting to 2D
    _2d_window = _1d_window.mm(_1d_window.t()).float().unsqueeze(0).unsqueeze(0)

    window = torch.Tensor(_2d_window.expand(channel, 1, window_size,
    ↪window_size).contiguous())

    ↪return window

```

```

[9]: # mean SSIM with SSIM applied locally over moving windows
# output = 1: the same image, output = 0 (or -1): very different
def mean_ssim(img1, img2, val_range, window_size=11, window=None,
    ↪size_average=True, full=False):

    L = val_range # L is the dynamic range of the pixel values (255 for 8-bit
    ↪grayscale images),

    pad = window_size // 2

    try:
        _, channels, height, width = img1.size()
    except:
        channels, height, width = img1.size()

    # if window is not provided, init one
    if window is None:
        real_size = min(window_size, height, width) # window should be atleast
    ↪11x11
        window = create_window(real_size, channel=channels).to(img1.device)

    # calculating the mu parameter (locally) for both images using a gaussian
    ↪filter
    # calculates the luminosity params
    mu1 = F.conv2d(img1, window, padding=pad, groups=channels)
    mu2 = F.conv2d(img2, window, padding=pad, groups=channels)

```

```

mu1_sq = mu1 ** 2
mu2_sq = mu2 ** 2
mu12 = mu1 * mu2

# now we calculate the sigma square parameter
# Sigma deals with the contrast component
sigma1_sq = F.conv2d(img1 * img1, window, padding=pad, groups=channels) -
↪mu1_sq
sigma2_sq = F.conv2d(img2 * img2, window, padding=pad, groups=channels) -
↪mu2_sq
sigma12 = F.conv2d(img1 * img2, window, padding=pad, groups=channels) -
↪mu12

# Some constants for stability
C1 = (0.01 ) ** 2 # NOTE: Removed L from here (ref PT implementation)
C2 = (0.03 ) ** 2

contrast_metric = (2.0 * sigma12 + C2) / (sigma1_sq + sigma2_sq + C2)
contrast_metric = torch.mean(contrast_metric)

numerator1 = 2 * mu12 + C1
numerator2 = 2 * sigma12 + C2
denominator1 = mu1_sq + mu2_sq + C1
denominator2 = sigma1_sq + sigma2_sq + C2

ssim_score = (numerator1 * numerator2) / (denominator1 * denominator2)

if size_average:
    ret = ssim_score.mean()
else:
    ret = ssim_score.mean(1).mean(1).mean(1)

if full:
    return ret, contrast_metric

return ret

```

```

[10]: # Helper functions to convert to Tensors
tensorify = lambda x: torch.Tensor(x.transpose((1, 0))).unsqueeze(0).float().
↪div(255.0)

```

```

[9]: # ### Example Usage ###

# img_path_temp = 'original/cameraman.png'
# img1 = io.imread(img_path_temp)
# img2 = io.imread(img_path_temp)

```

```

# # Check SSIM score of True image vs False Image
# _img1 = tensorify(img1)
# _img2 = tensorify(img2)
# true_vs_false = mean_ssim(_img1, _img2, val_range=255)
# print("True vs False Image SSIM Score:", true_vs_false)

```

## 0.0.6 SIREN Network Architecture

```

[11]: class SineLayer(nn.Module):
    # See paper sec. 3.2, final paragraph, and supplement Sec. 1.5 for
    ↪ discussion of omega_0.

    # If is_first=True, omega_0 is a frequency factor which simply multiplies
    ↪ the activations before the
    # nonlinearity. Different signals may require different omega_0 in the
    ↪ first layer - this is a
    # hyperparameter.

    # If is_first=False, then the weights will be divided by omega_0 so as to
    ↪ keep the magnitude of
    # activations constant, but boost gradients to the weight matrix (see
    ↪ supplement Sec. 1.5)

    def __init__(self, in_features, out_features, bias=True,
                  is_first=False, omega_0=30):
        super().__init__()
        self.omega_0 = omega_0
        self.is_first = is_first

        self.in_features = in_features
        self.linear = nn.Linear(in_features, out_features, bias=bias)

        self.init_weights()

    def init_weights(self):
        with torch.no_grad():
            if self.is_first:
                self.linear.weight.uniform_(-1 / self.in_features,
                                             1 / self.in_features)
            else:
                self.linear.weight.uniform_(-np.sqrt(6 / self.in_features) /
    ↪ self.omega_0,
                                             np.sqrt(6 / self.in_features) /
    ↪ self.omega_0)

    def forward(self, input):

```



```

        return torch.sin(self.omega_0 * self.linear(input))

    def forward_with_intermediate(self, input):
        # For visualization of activation distributions
        intermediate = self.omega_0 * self.linear(input)
        return torch.sin(intermediate), intermediate

class Siren(nn.Module):
    def __init__(self, in_features, hidden_features, hidden_layers,
        ↪out_features, outermost_linear=False,
            first_omega_0=30, hidden_omega_0=30.):
        super().__init__()

        self.net = []
        self.net.append(SineLayer(in_features, hidden_features,
                                   is_first=True, omega_0=first_omega_0))

        for i in range(hidden_layers):
            self.net.append(SineLayer(hidden_features, hidden_features,
                                       is_first=False, omega_0=hidden_omega_0))

        if outermost_linear:
            final_linear = nn.Linear(hidden_features, out_features)

            with torch.no_grad():
                final_linear.weight.uniform_(-np.sqrt(6 / hidden_features) /
        ↪hidden_omega_0,
                                                np.sqrt(6 / hidden_features) /
        ↪hidden_omega_0)

            self.net.append(final_linear)
        else:
            self.net.append(SineLayer(hidden_features, out_features,
                                       is_first=False, omega_0=hidden_omega_0))

        self.net = nn.Sequential(*self.net)

    def forward(self, coords):
        coords = coords.clone().detach().requires_grad_(True) # allows to take
        ↪derivative w.r.t. input
        output = self.net(coords)
        return output, coords

    def forward_with_activations(self, coords, retain_grad=False):
        '''Returns not only model output, but also intermediate activations.
        Only used for visualizing activations later!'''

```

```

        activations = OrderedDict()

        activation_count = 0
        x = coords.clone().detach().requires_grad_(True)
        activations['input'] = x
        for i, layer in enumerate(self.net):
            if isinstance(layer, SineLayer):
                x, intermed = layer.forward_with_intermediate(x)

                if retain_grad:
                    x.retain_grad()
                    intermed.retain_grad()

                activations['_'.join((str(layer.__class__), "%d" %_
↪activation_count)))] = intermed
                activation_count += 1
            else:
                x = layer(x)

                if retain_grad:
                    x.retain_grad()

                activations['_'.join((str(layer.__class__), "%d" %_
↪activation_count)))] = x
                activation_count += 1

        return activations

```

### 0.0.7 ELU Network Architecture

```

[12]: class ELULayer(nn.Module):

    def __init__(self, in_features, out_features, bias=True):
        super().__init__()

        self.in_features = in_features
        self.linear = nn.Linear(in_features, out_features, bias=bias)

        self.init_weights()

    def init_weights(self):
        with torch.no_grad():
            nn.init.xavier_uniform_(self.linear.weight)

    def forward(self, input):
        return F.elu(self.linear(input))

```

```

def forward_with_intermediate(self, input):
    # For visualization of activation distributions
    intermediate = self.linear(input)
    return F.elu(intermediate), intermediate

class Base(nn.Module):
    def __init__(self, in_features, hidden_features, hidden_layers,
        out_features, outermost_linear=False,
        first_omega_0=30, hidden_omega_0=30.):
        super().__init__()

        self.net = []
        self.net.append(ELULayer(in_features, hidden_features))

        for i in range(hidden_layers):
            self.net.append(ELULayer(hidden_features, hidden_features))

        if outermost_linear:
            final_linear = nn.Linear(hidden_features, out_features)

            with torch.no_grad():
                nn.init.xavier_uniform_(final_linear.weight)

            self.net.append(final_linear)
        else:
            self.net.append(ELULayer(hidden_features, out_features,
                is_first=False, omega_0=hidden_omega_0))

        self.net = nn.Sequential(*self.net)

    def forward(self, coords):
        coords = coords.clone().detach().requires_grad_(True) # allows to take
        derivative w.r.t. input
        output = self.net(coords)
        return output, coords

    def forward_with_activations(self, coords, retain_grad=False):
        '''Returns not only model output, but also intermediate activations.
        Only used for visualizing activations later!'''
        activations = OrderedDict()

        activation_count = 0
        x = coords.clone().detach().requires_grad_(True)
        activations['input'] = x
        for i, layer in enumerate(self.net):
            if isinstance(layer, SineLayer):

```

```

        x, intermed = layer.forward_with_intermediate(x)

        if retain_grad:
            x.retain_grad()
            intermed.retain_grad()

        activations['_'.join((str(layer.__class__), "%d" %
↪activation_count)))] = intermed
        activation_count += 1
    else:
        x = layer(x)

        if retain_grad:
            x.retain_grad()

        activations['_'.join((str(layer.__class__), "%d" %
↪activation_count)))] = x
        activation_count += 1

    return activations

```

### 0.0.8 Save video

```

[13]: import shutil
import subprocess

def output_video(net, img_path, niter, vidName='video_name.mp4'):

    image = ImageFitting(img_path=img_path, niter=niter)
    dataloader = DataLoader(image, batch_size=1, pin_memory=True, num_workers=0)

    net.cuda()

    if os.path.exists("tmp"):
        shutil.rmtree("tmp")
    os.makedirs("tmp")

    for step, batch in tqdm(enumerate(dataloader)):

        model_input = batch[0].cuda()
        gt = {key: value.cuda() for key, value in batch[1].items()}

        model_output, coords = net(model_input)
        img_grad = computeJacobianFull(coords, model_output, create_graph=True)
        img_laplacian = computeLaplaceFull(coords, img_grad, create_graph=False)

        fig, axes = plt.subplots(2,4, figsize=(18,6))

```

```

        axes[0,0].imshow(gt['pixels'][0].cpu().view(256,256).detach().numpy())
        axes[0,1].imshow(gt['grads'][0].norm(dim=-1).cpu().view(256,256).
↳detach().numpy())
        axes[0,2].imshow(gt['laplace'][0].cpu().view(256,256).detach().numpy())
        axes[0,3].imshow(gt['dIdt'][0].cpu().view(256,256).detach().numpy())
        axes[1,0].imshow(model_output[0].cpu().view(256,256).detach().numpy())
        axes[1,1].imshow(img_grad[0][:,-1].norm(dim=-1).cpu().view(256,256).
↳detach().numpy())
        axes[1,2].imshow(img_laplacian[0].cpu().view(256,256).detach().numpy())
        axes[1,3].imshow(img_grad[0][:,-1].cpu().view(256,256).detach().numpy())

    fig.savefig("tmp/file%02d.png" % step)

    subprocess.call([
        'ffmpeg', '-framerate', '2', '-i', 'tmp/file%02d.png', '-r', '30',
↳'-pix_fmt', 'yuv420p',
        vidName
    ])

    shutil.rmtree("tmp")

```

### 0.0.9 Train Network

```

[14]: def train(net, writer, img_path, niter, total_epochs=50, lr=[1e-4], beta_0=1,
↳beta_1=1, beta_2=1, beta_3=1,
        cyclic=False, decay_exp=False, decay_multi=False):

    """Args:
        net: Network to Train
        writer: SummaryWriter for logging
        img_path: path to default state image
        niter: number of steps to apply diffusion (0 means only 1 image)
        total_epochs: number of epochs to train
        beta_0: constant for loss on pixel value
        beta_1: constant for loss on gradients
        beta_2: constant for loss on laplacian
        beta_3: constant for loss on pixel time derivative
        cyclic: CyclicLearning rate (allows better learning)"""

    image = ImageFitting(img_path=img_path, niter=niter)
    dataloader = DataLoader(image, batch_size=1, pin_memory=True, num_workers=0)

    net.cuda()

    epochs_til_summary = 10 #UPDATE ACCORDINGLY

```

```

steps_til_summary = 5 #UPDATE ACCORDINGLY

optim = torch.optim.Adam(lr=lr[0], params=net.parameters())

if decay_multi:
    m = np.floor(total_epochs/4)
    scheduler = torch.optim.lr_scheduler.MultiStepLR(optim,
↪milestones=[m*1,m*2,m*3], gamma=0.1)

if decay_exp:
    scheduler = torch.optim.lr_scheduler.ExponentialLR(optim, gamma=0.9)

if cyclic:
    scheduler = torch.optim.lr_scheduler.CyclicLR(optim, base_lr=lr[1],
↪max_lr=lr[0], step_size_up=250, cycle_momentum=False)

print("-----Begin Training-----")
for epoch in range(1, total_epochs + 1):

    epoch_loss = 0.0
    epoch_pixel_loss = 0.0
    epoch_grad_loss = 0.0
    epoch_laplacian_loss = 0.0
    epoch_dIdt_loss = 0.0

    epoch_pixel_ssim = 0.0
    epoch_grad_ssim = 0.0
    epoch_laplacian_ssim = 0.0
    epoch_dIdt_ssim = 0.0

    for step, batch in tqdm(enumerate(dataloader)):

        model_input = batch[0].cuda()
        gt = {key: value.cuda() for key, value in batch[1].items()}

        model_output, coords = net(model_input)

        pixel_loss, grad_loss, laplacian_loss, dIdt_loss, pixel_ssim,
↪grad_ssim, laplacian_ssim, dIdt_ssim = calcLoss(coords, model_output, gt)

        loss = beta_0 * pixel_loss + beta_1 * grad_loss + beta_2 *
↪laplacian_loss + beta_3 * dIdt_loss

        epoch_loss += model_output.shape[0] * loss.item()
        epoch_pixel_loss += model_output.shape[0] * pixel_loss.item()
        epoch_grad_loss += model_output.shape[0] * grad_loss.item()

```

```

epoch_laplacian_loss += model_output.shape[0] * laplacian_loss.
↪item()
epoch_dIdt_loss += model_output.shape[0] * dIdt_loss.item()

epoch_pixel_ssim += model_output.shape[0] * pixel_ssim.item()
epoch_grad_ssim += model_output.shape[0] * grad_ssim.item()
epoch_laplacian_ssim += model_output.shape[0] * laplacian_ssim.
↪item()
epoch_dIdt_ssim += model_output.shape[0] * dIdt_ssim.item()

if not epoch % epochs_til_summary and step % steps_til_summary == 0:
↪steps_til_summary - 1:

    pixel_output = model_output[0].view(1, -1, 256, 256)
    pixel_gt = gt['pixels'][0].view(1, -1, 256, 256)
    img_grid_pixel = torchvision.utils.make_grid(torch.
↪cat((pixel_gt, pixel_output), 0), 2)
    img_grid_pixel = img_grid_pixel * 0.5 + 0.5
    writer.add_image('pixels', img_grid_pixel, epoch * 0.5
↪len(dataloader) + step + 1)

    img_grad = computeJacobianFull(coords, model_output, 0.5
↪create_graph=True)
    grad_output = img_grad[0,:,:-1].norm(dim=-1).view(1, -1, 256, 0.5
↪256)
    grad_gt = gt['grads'][0].norm(dim=-1).view(1, -1, 256, 256)
    img_grid_grad = torchvision.utils.make_grid(torch.cat((grad_gt, 0.5
↪grad_output), 0), 2)
    writer.add_image('grads', img_grid_grad, epoch * 0.5
↪len(dataloader) + step + 1)

    img_laplacian = computeLaplaceFull(coords, img_grad, 0.5
↪create_graph=False)
    laplacian_output = img_laplacian[0].view(1, -1, 256, 256)
    laplacian_gt = gt['laplace'][0].view(1, -1, 256, 256)
    img_grid_laplacian = torchvision.utils.make_grid(torch.
↪cat((laplacian_gt, laplacian_output), 0), 2)
    writer.add_image('laplacians', img_grid_laplacian, epoch * 0.5
↪len(dataloader) + step + 1)

    dIdt_output = img_grad[0,:,:-1].view(1, -1, 256, 256)
    dIdt_gt = gt['dIdt'][0].view(1, -1, 256, 256)
    img_grid_dIdt = torchvision.utils.make_grid(torch.cat((dIdt_gt, 0.5
↪dIdt_output), 0), 2)

```

```

        writer.add_image('dIdt', img_grid_dIdt, epoch * len(dataloader)
↪ + step + 1)

        # fig, axes = plt.subplots(2,4, figsize=(18,6))
        # axes[0,0].imshow(gt['pixels'][0].cpu().view(256,256).detach().
↪ numpy())
        # axes[0,1].imshow(gt['grads'][0].norm(dim=-1).cpu().
↪ view(256,256).detach().numpy())
        # axes[0,2].imshow(gt['laplace'][0].cpu().view(256,256).
↪ detach().numpy())
        # axes[0,3].imshow(gt['dIdt'][0].cpu().view(256,256).detach().
↪ numpy())
        # axes[1,0].imshow(model_output[0].cpu().view(256,256).detach().
↪ numpy())
        # axes[1,1].imshow(img_grad[0][:,-1].norm(dim=-1).cpu().
↪ view(256,256).detach().numpy())
        # axes[1,2].imshow(img_laplacian[0].cpu().view(256,256).
↪ detach().numpy())
        # axes[1,3].imshow(img_grad[0][:,-1].cpu().view(256,256).
↪ detach().numpy())
        # plt.show()

        optim.zero_grad()
        loss.backward()
        optim.step()

        if cyclic or decay_exp or decay_multi:
            scheduler.step()

        # logging epoch loss
        writer.add_scalar('epoch_loss/total', epoch_loss/len(image), epoch)
        writer.add_scalar('epoch_loss/pixel', epoch_pixel_loss/len(image),
↪ epoch)
        writer.add_scalar('epoch_loss/grad', epoch_grad_loss/len(image), epoch)
        writer.add_scalar('epoch_loss/laplacian', epoch_laplacian_loss/
↪ len(image), epoch)
        writer.add_scalar('epoch_loss/dIdt', epoch_dIdt_loss/len(image), epoch)
        print("Epoch %d, Epoch loss: total %.6f, pixel %.6f, grad %.6f,
↪ laplacian %.6f, dIdt %.6f" % (epoch, epoch_loss/len(image),
↪ epoch_pixel_loss/len(image), epoch_grad_loss/len(image),
↪ epoch_laplacian_loss/len(image), epoch_dIdt_loss/len(image)))

        # logging ssim loss
        writer.add_scalar('epoch_ssim/pixel', epoch_pixel_ssim/len(image),
↪ epoch)
        writer.add_scalar('epoch_ssim/grad', epoch_grad_ssim/len(image), epoch)

```



```

        writer.add_scalar('epoch_ssim/laplacian', epoch_laplacian_ssim/
↪len(image), epoch)
        writer.add_scalar('epoch_ssim/dIdt', epoch_dIdt_ssim/len(image), epoch)
        print("Epoch %d, Epoch SSIM: pixel %0.6f, grad %0.6f, laplacian %0.6f, ↪
↪dIdt %0.6f" % (epoch, epoch_pixel_ssim/len(image), epoch_grad_ssim/
↪len(image), epoch_laplacian_ssim/len(image), epoch_dIdt_ssim/len(image)))

writer.add_graph(net, model_input)
print("-----Finished-----")

```

## 0.0.10 Baselines & Activation Experiments

### SIREN Baselines

```
[ ]: torch.cuda.empty_cache()
```

```
[ ]: total_epochs = 20
```

```
[ ]: # SIREN, learn only with the observed pixel values
writer = SummaryWriter('runs/siren/cameraman_experiment_pixels')

img_siren = Siren(in_features=3, out_features=1, hidden_features=512,
                  hidden_layers=3, outermost_linear=True)

train(img_siren, writer, img_path='original/cameraman.png', niter=10, ↪
↪total_epochs=total_epochs, lr=[1e-4],
      beta_0=1, beta_1=0, beta_2=0, beta_3=0)

writer.close()

output_video(img_siren, img_path='original/cameraman.png', niter=10, ↪
↪vidName='videos/siren/cameraman_experiment_pixels'+ '_video.mp4')
```

```
[ ]: # SIREN, learn only with the observed jacobians (first derivative in space)
writer = SummaryWriter('runs/siren/cameraman_experiment_grads')

img_siren = Siren(in_features=3, out_features=1, hidden_features=512,
                  hidden_layers=3, outermost_linear=True)

train(img_siren, writer, img_path='original/cameraman.png', niter=1, ↪
↪total_epochs=total_epochs, lr=[1e-4],
      beta_0=0, beta_1=1, beta_2=0, beta_3=0)

writer.close()
```

```
output_video(img_siren, img_path='original/cameraman.png', niter=10,
↳vidName='videos/siren/cameraman_experiment_grads'+ '_video.mp4')
```

```
[ ]: # SIREN, learns only with the observed laplacians (2nd derivative in space)
writer = SummaryWriter('runs/siren/cameraman_experiment_laplace')

img_siren = Siren(in_features=3, out_features=1, hidden_features=512,
                  hidden_layers=3, outermost_linear=True)

train(img_siren, writer, img_path='original/cameraman.png', niter=1,
↳total_epochs=total_epochs, lr=[1e-4],
    beta_0=0, beta_1=0, beta_2=1, beta_3=0)

writer.close()

output_video(img_siren, img_path='original/cameraman.png', niter=10,
↳vidName='videos/siren/cameraman_experiment_laplace'+ '_video.mp4')
```

```
[ ]: # SIREN, learn only with the observed derivative in time (3rd derivative)
writer = SummaryWriter('runs/siren/cameraman_experiment_dIdt')

img_siren = Siren(in_features=3, out_features=1, hidden_features=512,
                  hidden_layers=3, outermost_linear=True)

train(img_siren, writer, img_path='original/cameraman.png', niter=1,
↳total_epochs=total_epochs, lr=[1e-4],
    beta_0=0, beta_1=0, beta_2=0, beta_3=1)

writer.close()

output_video(img_siren, img_path='original/cameraman.png', niter=10,
↳vidName='videos/siren/cameraman_experiment_dIdt'+ '_video.mp4')
```

```
[ ]: # SIREN, learn with all data, equally weighted
writer = SummaryWriter('runs/siren/cameraman_experiment_all')

img_siren = Siren(in_features=3, out_features=1, hidden_features=512,
                  hidden_layers=3, outermost_linear=True)

train(img_siren, writer, img_path='original/cameraman.png', niter=1,
↳total_epochs=total_epochs, lr=[1e-4],
    beta_0=1, beta_1=1, beta_2=1, beta_3=1)

writer.close()

output_video(img_siren, img_path='original/cameraman.png', niter=10,
↳vidName='videos/siren/cameraman_experiment_all'+ '_video.mp4')
```

## Elu Baselines

```
[ ]: # Base, learn only with the observed pixel values
writer = SummaryWriter('runs/base/cameraman_experiment_pixels')

img_base = Base(in_features=3, out_features=1, hidden_features=512,
                hidden_layers=3, outermost_linear=True)

train(img_base, writer, img_path='original/cameraman.png', niter=1,
      ↪total_epochs=total_epochs, lr=[1e-4],
      beta_0=1, beta_1=0, beta_2=0, beta_3=0)

writer.close()
output_video(img_siren, img_path='original/cameraman.png', niter=10,
      ↪vidName='videos/base/cameraman_experiment_pixels'+ '_video.mp4')

[ ]: # Base, learn only with the observed jacobians (first derivative in space)
writer = SummaryWriter('runs/base/cameraman_experiment_grads')

img_base = Base(in_features=3, out_features=1, hidden_features=512,
                hidden_layers=3, outermost_linear=True)

train(img_base, writer, img_path='original/cameraman.png', niter=1,
      ↪total_epochs=total_epochs, lr=[1e-4],
      beta_0=0, beta_1=1, beta_2=0, beta_3=0)

writer.close()

output_video(img_siren, img_path='original/cameraman.png', niter=10,
      ↪vidName='videos/base/cameraman_experiment_grads'+ '_video.mp4')

[ ]: # Base, learn only with the observed laplacians (2nd derivative in space)
writer = SummaryWriter('runs/base/cameraman_experiment_laplace')

img_base = Base(in_features=3, out_features=1, hidden_features=512,
                hidden_layers=3, outermost_linear=True)

train(img_base, writer, img_path='original/cameraman.png', niter=1,
      ↪total_epochs=total_epochs, lr=[1e-4],
      beta_0=0, beta_1=0, beta_2=1, beta_3=0)

writer.close()
output_video(img_siren, img_path='original/cameraman.png', niter=10,
      ↪vidName='videos/base/cameraman_experiment_laplace'+ '_video.mp4')

[ ]: # Base, learn only with the observed derivative in time (3rd derivative)
writer = SummaryWriter('runs/base/cameraman_experiment_d1dt')
```

```

img_base = Base(in_features=3, out_features=1, hidden_features=512,
                hidden_layers=3, outermost_linear=True)

train(img_base, writer, img_path='original/cameraman.png', niter=1,
      ↪total_epochs=total_epochs, lr=[1e-4],
        beta_0=0, beta_1=0, beta_2=0, beta_3=1)

writer.close()

output_video(img_siren, img_path='original/cameraman.png', niter=10,
  ↪vidName='videos/base/cameraman_experiment_d1dt' + '_video.mp4')

```

```

[ ]: # Base, learn with all data, equally weighted
writer = SummaryWriter('runs/base/cameraman_experiment_all')

img_base = Base(in_features=3, out_features=1, hidden_features=512,
                hidden_layers=3, outermost_linear=True)

train(img_base, writer, img_path='original/cameraman.png', niter=1,
      ↪total_epochs=total_epochs, lr=[1e-4],
        beta_0=1, beta_1=1, beta_2=1, beta_3=1)

writer.close()

output_video(img_siren, img_path='original/cameraman.png', niter=10,
  ↪vidName='videos/base/cameraman_experiment_all' + '_video.mp4')

```

## 0.0.11 Experiments

### Experiments Helper Functions

```

[ ]: # @min_beta_sum: minimum sum of all four beta values
# @return: a list of 4 beta values, summing to at least min_beta_sum
def generate_random_beta_combos(min_beta_sum=0.1):
    possible_values = [1.0, 0.1, 0.01, 0.001, 0.0]
    betas = [0, 0, 0, 0]

    while np.sum(betas) <= min_beta_sum:
        betas = [random.choice(possible_values) for i in range(4)]

    return betas

# @betas: a list of beta values
# @return: a string with '-' between all beta values
def b_to_string(betas):
    return '_'.join(map(str, betas))

```

```
[ ]: # runs one experiment with the elu activation function
def run_elu(model_path, betas, total_epochs, lr, cyclic=False, decay_exp=False,
    ↪decay_multi=False):
    writer = SummaryWriter(model_path)

    img_base = Base(in_features=3, out_features=1, hidden_features=512,
        hidden_layers=3, outermost_linear=True)

    train(img_base, writer, img_path='original/cameraman.png', niter=10,
        total_epochs=total_epochs, lr=lr,
        beta_0=betas[0], beta_1=betas[1], beta_2=betas[2], beta_3=betas[3],
        cyclic=cyclic, decay_exp=decay_exp, decay_multi=decay_multi)

    writer.close()

    output_video(img_base, img_path='original/cameraman.png', niter=10,
    ↪vidName='/home/jupyter/videos/' + model_path+ '_video.mp4')
```

```
[ ]: # runs one experiment with the SIREN (periodic) activation function
def run_siren(model_path, betas, total_epochs, lr, cyclic=False,
    ↪decay_exp=False, decay_multi=False):
    writer = SummaryWriter(model_path)

    img_siren = Siren(in_features=3, out_features=1, hidden_features=512,
        hidden_layers=3, outermost_linear=True)

    train(img_siren, writer, img_path='original/cameraman.png', niter=10,
        total_epochs=total_epochs, lr=lr,
        beta_0=betas[0], beta_1=betas[1], beta_2=betas[2], beta_3=betas[3],
        cyclic=cyclic, decay_exp=decay_exp, decay_multi=decay_multi)

    writer.close()

    output_video(img_siren, img_path='original/cameraman.png', niter=10,
    ↪vidName='/home/jupyter/videos/' + model_path+ '_video.mp4')
```

## Run the experiments

```
[25]: # settings
model_path = 'runs/cameraman/experiments'
total_epochs = 20 #100
#learning_rates = [1e-4, 1e-5, 1e-6, 1e-7]
learning_rates = [1e-4]

# keep conducting experiments until we've reached the desired amount
num_experiments = 0
while num_experiments < 100:
```

```

torch.cuda.empty_cache()

# get a random combination of betas
betas = generate_random_beta_combos()
model_path_b = model_path + "/" + b_to_string(betas)

if Path(model_path_b).exists():
    continue

os.mkdir('/home/jupyter/videos/' + model_path_b)

### Part A. Models with elu activation ###
#     model_path_act = model_path + '/elu_'

#     ## learning rate experiments ##

#     # 1. run with uniform learning rates
#     for uniform_lr in learning_rates:
#         model_path_full = model_path_act + 'uniformlr_' + "{:.0e}".
#         ↪format(uniform_lr)
#         run_elu(model_path_full, betas, total_epochs, [uniform_lr])

#     # 2. run with decaying learning rates
#     initial_lr = learning_rates[0]

#     # 2.1 multi-step: decay_multi = True
#     model_path_full = model_path_act + '_decay_multi_' + "{:.0e}".
#     ↪format(initial_lr)
#     run_elu(model_path_full, betas, total_epochs, [initial_lr], ↪
#     ↪decay_multi=True)

#     # 2.2 exponential: decay_exp = True
#     model_path_full = model_path_act + '_decay_exp_' + "{:.0e}".
#     ↪format(initial_lr)
#     run_elu(model_path_full, betas, total_epochs, [initial_lr], ↪
#     ↪decay_exp=True)

#     # 3. run with cyclic learning rate
#     max_lr = learning_rates[0]
#     min_lr = learning_rates[-1]
#     model_path_full = model_path_act + '_cyclic_' + "{:.0e}".format(max_lr) + ↪
#     ↪"_" + "{:.0e}".format(min_lr)
#     run_elu(model_path_full, betas, total_epochs, [max_lr, min_lr], ↪
#     ↪cyclic=True)

```

```

### Part B. Models with SIREN (periodic) activation ###
model_path_act = model_path_b + '/siren_'

## learning rate experiments ##

# 1. run with uniform learning rates
for uniform_lr in learning_rates:
    model_path_full = model_path_act + 'uniformlr_' + "{:.0e}".
    ↪format(uniform_lr)
    run_siren(model_path_full, betas, total_epochs, [uniform_lr])

# 2. run with decaying learning rates
#     initial_lr = learning_rates[0]

#     # 2.1 multi-step: decay_multi = True
#     model_path_full = model_path_act + '_decay_multi_' + "{:.0e}".
    ↪format(initial_lr)
#     run_siren(model_path_full, betas, total_epochs, [initial_lr], ↪
    ↪decay_multi=True)

#     # 2.2 exponential: decay_exp = True
#     model_path_full = model_path_act + '_decay_exp_' + "{:.0e}".
    ↪format(initial_lr)
#     run_siren(model_path_full, betas, total_epochs, [initial_lr], ↪
    ↪decay_exp=True)

#     # 3. run with cyclic learning rate
#     max_lr = learning_rates[0]
#     min_lr = learning_rates[-1]
#     model_path_full = model_path_act + '_cyclic_' + "{:.0e}".format(max_lr) + ↪
    ↪"_" + "{:.0e}".format(min_lr)
#     run_siren(model_path_full, betas, total_epochs, [max_lr, min_lr], ↪
    ↪cyclic=True)

# finished one more experiment
num_experiments += 1
print("finished experiment #", num_experiments)

```

-----Generating Data-----

100%| | 10/10 [00:00<00:00, 1167.06it/s]

-----Finished-----

-----Begin Training-----

10it [00:25, 2.54s/it]

Epoch 1, Epoch loss: total 1792.176624, pixel 0.341292, grad 8.266688, laplacian 132.000001, dIdt 17916.950391

Epoch 1, Epoch SSIM: pixel 0.065567, grad 0.027473, laplacian 0.000044, dIdt 0.000925  
10it [00:26, 2.60s/it]

Epoch 2, Epoch loss: total 1779.828564, pixel 0.397523, grad 10.770234, laplacian 1250.852451, dIdt 17781.694043  
Epoch 2, Epoch SSIM: pixel 0.002419, grad 0.012909, laplacian 0.000003, dIdt 0.000297  
10it [00:25, 2.58s/it]

Epoch 3, Epoch loss: total 1968.496521, pixel 0.328358, grad 19.897638, laplacian 170625.655823, dIdt 17975.225586  
Epoch 3, Epoch SSIM: pixel 0.052166, grad 0.013906, laplacian 0.000006, dIdt 0.000214  
10it [00:26, 2.63s/it]

Epoch 4, Epoch loss: total 2140.969055, pixel 0.339396, grad 24.207743, laplacian 341306.188379, dIdt 17992.992383  
Epoch 4, Epoch SSIM: pixel 0.022648, grad 0.010504, laplacian -0.000002, dIdt 0.000158  
10it [00:25, 2.57s/it]

Epoch 5, Epoch loss: total 2082.272510, pixel 0.336864, grad 24.372157, laplacian 285695.541211, dIdt 17962.156738  
Epoch 5, Epoch SSIM: pixel -0.005028, grad 0.013174, laplacian 0.000001, dIdt -0.000004  
10it [00:26, 2.64s/it]

Epoch 6, Epoch loss: total 2030.863867, pixel 0.330846, grad 21.713527, laplacian 238419.892285, dIdt 17920.913770  
Epoch 6, Epoch SSIM: pixel 0.020352, grad 0.010625, laplacian 0.000001, dIdt 0.000051  
10it [00:25, 2.55s/it]

Epoch 7, Epoch loss: total 2007.651086, pixel 0.335302, grad 22.325742, laplacian 219646.560547, dIdt 17876.468848  
Epoch 7, Epoch SSIM: pixel 0.038611, grad 0.011424, laplacian 0.000001, dIdt 0.000057  
10it [00:25, 2.58s/it]

Epoch 8, Epoch loss: total 2051.837268, pixel 0.340840, grad 23.098742, laplacian 264609.363672, dIdt 17868.639453  
Epoch 8, Epoch SSIM: pixel 0.055094, grad 0.010403, laplacian 0.000001, dIdt 0.000050  
10it [00:26, 2.64s/it]

Epoch 9, Epoch loss: total 2079.527124, pixel 0.345991, grad 27.621856, laplacian 297686.437500, dIdt 17814.670801



Epoch 9, Epoch SSIM: pixel 0.068804, grad 0.006542, laplacian 0.000000, dIdt  
-0.000033  
10it [00:28, 2.84s/it]

Epoch 10, Epoch loss: total 2043.798621, pixel 0.377362, grad 32.715271,  
laplacian 263866.433984, dIdt 17795.220508  
Epoch 10, Epoch SSIM: pixel 0.088739, grad 0.006923, laplacian 0.000001, dIdt  
-0.000082  
10it [00:26, 2.62s/it]

Epoch 11, Epoch loss: total 2009.568262, pixel 0.392238, grad 34.392525,  
laplacian 224966.730469, dIdt 17841.748730  
Epoch 11, Epoch SSIM: pixel 0.074478, grad 0.007866, laplacian -0.000000, dIdt  
-0.000023  
10it [00:25, 2.54s/it]

Epoch 12, Epoch loss: total 2110.206665, pixel 0.402601, grad 34.824779,  
laplacian 326729.953516, dIdt 17830.392871  
Epoch 12, Epoch SSIM: pixel 0.070939, grad 0.005185, laplacian 0.000000, dIdt  
-0.000122  
10it [00:25, 2.59s/it]

Epoch 13, Epoch loss: total 2138.444836, pixel 0.428828, grad 35.287197,  
laplacian 360075.812891, dIdt 17779.048828  
Epoch 13, Epoch SSIM: pixel 0.108333, grad 0.006109, laplacian 0.000000, dIdt  
-0.000063  
10it [00:25, 2.54s/it]

Epoch 14, Epoch loss: total 2212.184265, pixel 0.497210, grad 52.824970,  
laplacian 444346.478906, dIdt 17672.876563  
Epoch 14, Epoch SSIM: pixel 0.093855, grad 0.005503, laplacian -0.000000, dIdt  
-0.000104  
10it [00:25, 2.56s/it]

Epoch 15, Epoch loss: total 2292.813879, pixel 0.589648, grad 58.364088,  
laplacian 524507.364844, dIdt 17676.584180  
Epoch 15, Epoch SSIM: pixel 0.107817, grad 0.006651, laplacian 0.000000, dIdt  
-0.000115  
10it [00:26, 2.65s/it]

Epoch 16, Epoch loss: total 2330.021143, pixel 0.584426, grad 66.697298,  
laplacian 560413.314063, dIdt 17689.566211  
Epoch 16, Epoch SSIM: pixel 0.103204, grad 0.005465, laplacian 0.000000, dIdt  
-0.000078  
10it [00:25, 2.53s/it]

Epoch 17, Epoch loss: total 2218.156934, pixel 0.763054, grad 64.608762,  
laplacian 464656.876563, dIdt 17526.723535

Epoch 17, Epoch SSIM: pixel 0.106037, grad 0.005756, laplacian -0.000000, dIdt 0.000014

10it [00:26, 2.61s/it]

Epoch 18, Epoch loss: total 2318.970398, pixel 0.832172, grad 89.430962, laplacian 560110.987500, dIdt 17579.377637

Epoch 18, Epoch SSIM: pixel 0.069979, grad 0.003898, laplacian 0.000000, dIdt 0.000029

10it [00:25, 2.55s/it]

Epoch 19, Epoch loss: total 2567.778101, pixel 0.859720, grad 104.167003, laplacian 810264.751563, dIdt 17565.493945

Epoch 19, Epoch SSIM: pixel 0.085459, grad 0.005708, laplacian 0.000000, dIdt -0.000080

10it [00:28, 2.87s/it]

Epoch 20, Epoch loss: total 2467.536633, pixel 0.888009, grad 109.146054, laplacian 705958.093750, dIdt 17605.813672

Epoch 20, Epoch SSIM: pixel 0.074682, grad 0.002666, laplacian 0.000000, dIdt -0.000151

-----Finished-----

-----Generating Data-----

100%| | 10/10 [00:00<00:00, 1255.82it/s]

-----Finished-----

10it [00:15, 1.57s/it]

ffmpeg version 4.1.9-0+deb10u1 Copyright (c) 2000-2022 the FFmpeg developers  
built with gcc 8 (Debian 8.3.0-6)

configuration: --prefix=/usr --extra-version=0+deb10u1 --toolchain=hardened  
--libdir=/usr/lib/x86\_64-linux-gnu --incdir=/usr/include/x86\_64-linux-gnu  
--arch=amd64 --enable-gpl --disable-stripping --enable-avresample --disable-  
filter=resample --enable-avisynth --enable-gnutls --enable-ladspa --enable-  
libaom --enable-libass --enable-libbluray --enable-libbs2b --enable-libcaca  
--enable-libcdio --enable-libcodec2 --enable-libflite --enable-libfontconfig  
--enable-libfreetype --enable-libfribidi --enable-libgme --enable-libgsm  
--enable-libjack --enable-libmp3lame --enable-libmysofa --enable-libopenjpeg  
--enable-libopenmpt --enable-libopus --enable-libpulse --enable-librsvg  
--enable-librubberband --enable-libshine --enable-libsnappy --enable-libsoxr  
--enable-libspeex --enable-libssh --enable-libtheora --enable-libtwolame  
--enable-libvidstab --enable-libvorbis --enable-libvpx --enable-libwavpack  
--enable-libwebp --enable-libx265 --enable-libxml2 --enable-libxvid --enable-  
libzmq --enable-libzvbi --enable-lv2 --enable-omx --enable-openal --enable-  
opengl --enable-sdl2 --enable-libdc1394 --enable-libdrm --enable-libiec61883  
--enable-chromaprint --enable-frei0r --enable-libx264 --enable-shared

libavutil 56. 22.100 / 56. 22.100

libavcodec 58. 35.100 / 58. 35.100

libavformat 58. 20.100 / 58. 20.100

```

libavdevice      58.  5.100 / 58.  5.100
libavfilter      7. 40.101 /  7. 40.101
libavresample    4.  0.  0 /  4.  0.  0
libswscale       5.  3.100 /  5.  3.100
libswresample    3.  3.100 /  3.  3.100
libpostproc     55.  3.100 / 55.  3.100
Input #0, image2, from 'tmp/file%02d.png':
  Duration: 00:00:05.00, start: 0.000000, bitrate: N/A
    Stream #0:0: Video: png, rgba(pc), 1296x432 [SAR 2835:2835 DAR 3:1], 2 fps,
2 tbr, 2 tbn, 2 tbc
Stream mapping:
  Stream #0:0 -> #0:0 (png (native) -> h264 (libx264))
Press [q] to stop, [?] for help
[libx264 @ 0x563c29a93f00] using SAR=1/1
[libx264 @ 0x563c29a93f00] using cpu capabilities: MMX2 SSE2Fast SSSE3 SSE4.2
AVX FMA3 BMI2 AVX2
[libx264 @ 0x563c29a93f00] profile High, level 3.1
[libx264 @ 0x563c29a93f00] 264 - core 155 r2917 0a84d98 - H.264/MPEG-4 AVC codec
- Copyleft 2003-2018 - http://www.videolan.org/x264.html - options: cabac=1
ref=3 deblock=1:0:0 analyse=0x3:0x113 me=hex subme=7 psy=1 psy_rd=1.00:0.00
mixed_ref=1 me_range=16 chroma_me=1 trellis=1 8x8dct=1 cqm=0 deadzone=21,11
fast_pskip=1 chroma_qp_offset=-2 threads=3 lookahead_threads=1 sliced_threads=0
nr=0 decimate=1 interlaced=0 bluray_compat=0 constrained_intra=0 bframes=3
b_pyramid=2 b_adapt=1 b_bias=0 direct=1 weightb=1 open_gop=0 weightp=2
keyint=250 keyint_min=25 scenecut=40 intra_refresh=0 rc_lookahead=40 rc=crf
mbtree=1 crf=23.0 qcomp=0.60 qpmin=0 qpmax=69 qpstep=4 ip_ratio=1.40 aq=1:1.00
Output #0, mp4, to '/home/jupyter/videos/runs/cameraman/experiments/1.0_0.001_0.
001_0.1/siren_uniformlr_1e-04_video.mp4':
  Metadata:
    encoder      : Lavf58.20.100
    Stream #0:0: Video: h264 (libx264) (avc1 / 0x31637661), yuv420p, 1296x432
[SAR 1:1 DAR 3:1], q=-1--1, 30 fps, 15360 tbn, 30 tbc
    Metadata:
      encoder      : Lavc58.35.100 libx264
    Side data:
      cpb: bitrate max/min/avg: 0/0/0 buffer size: 0 vbv_delay: -1
frame= 150 fps=101 q=-1.0 Lsize=      259kB time=00:00:04.90 bitrate=
433.6kbts/s dup=140 drop=0 speed= 3.3x
video:257kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB muxing
overhead: 1.008557%
[libx264 @ 0x563c29a93f00] frame I:1      Avg QP:15.97  size: 51228
[libx264 @ 0x563c29a93f00] frame P:38     Avg QP:19.11  size:  5304
[libx264 @ 0x563c29a93f00] frame B:111    Avg QP:14.80  size:    85
[libx264 @ 0x563c29a93f00] consecutive B-frames:  1.3%  0.0%  0.0% 98.7%
[libx264 @ 0x563c29a93f00] mb I  I16..4: 51.8% 23.2% 25.0%
[libx264 @ 0x563c29a93f00] mb P  I16..4:  0.6%  2.0%  1.0%  P16..4:  3.8%  0.9%
1.0%  0.0%  0.0%  skip:90.6%
[libx264 @ 0x563c29a93f00] mb B  I16..4:  0.1%  0.0%  0.0%  B16..8:  2.9%  0.0%

```

```

0.0% direct: 0.0% skip:96.9% L0:47.4% L1:52.5% BI: 0.1%
[libx264 @ 0x563c29a93f00] 8x8 transform intra:41.5% inter:60.4%
[libx264 @ 0x563c29a93f00] coded y,uvDC,uvAC intra: 56.7% 61.7% 58.2% inter:
0.8% 1.1% 0.6%
[libx264 @ 0x563c29a93f00] i16 v,h,dc,p: 68% 25% 7% 1%
[libx264 @ 0x563c29a93f00] i8 v,h,dc,ddl,ddr,vr,hd,vl,hu: 17% 13% 20% 8% 7%
8% 8% 9% 12%
[libx264 @ 0x563c29a93f00] i4 v,h,dc,ddl,ddr,vr,hd,vl,hu: 23% 21% 15% 7% 6%
7% 7% 7% 7%
[libx264 @ 0x563c29a93f00] i8c dc,h,v,p: 57% 19% 17% 7%
[libx264 @ 0x563c29a93f00] Weighted P-Frames: Y:0.0% UV:0.0%
[libx264 @ 0x563c29a93f00] ref P L0: 73.9% 18.9% 6.5% 0.7%
[libx264 @ 0x563c29a93f00] ref B L0: 69.1% 30.5% 0.4%
[libx264 @ 0x563c29a93f00] ref B L1: 97.1% 2.9%
[libx264 @ 0x563c29a93f00] kb/s:419.62

finished experiment # 1
-----Generating Data-----

100%|      | 10/10 [00:00<00:00, 1172.08it/s]

-----Finished-----
-----Begin Training-----

10it [00:25, 2.52s/it]

Epoch 1, Epoch loss: total 18.795853, pixel 0.293181, grad 7.466433, laplacian
113.264881, dIdt 17918.147949
Epoch 1, Epoch SSIM: pixel 0.024657, grad 0.034233, laplacian 0.000022, dIdt
0.001166

10it [00:26, 2.60s/it]

Epoch 2, Epoch loss: total 56.686100, pixel 0.270249, grad 6.993146, laplacian
496.902505, dIdt 17916.885645
Epoch 2, Epoch SSIM: pixel 0.099422, grad 0.036360, laplacian 0.000010, dIdt
0.002163

10it [00:25, 2.52s/it]

Epoch 3, Epoch loss: total 114.659120, pixel 0.272526, grad 7.131689, laplacian
1075.247021, dIdt 17926.952930
Epoch 3, Epoch SSIM: pixel 0.103362, grad 0.041289, laplacian 0.000006, dIdt
0.002324

10it [00:25, 2.59s/it]

Epoch 4, Epoch loss: total 66.207528, pixel 0.288026, grad 7.578935, laplacian
586.257114, dIdt 17923.577637
Epoch 4, Epoch SSIM: pixel 0.129646, grad 0.033112, laplacian 0.000006, dIdt
0.003137

10it [00:25, 2.55s/it]

```

Epoch 5, Epoch loss: total 99.082601, pixel 0.266842, grad 6.913469, laplacian 921.664624, dIdt 17924.849121  
Epoch 5, Epoch SSIM: pixel 0.166034, grad 0.045877, laplacian 0.000004, dIdt 0.001998  
10it [00:26, 2.61s/it]

Epoch 6, Epoch loss: total 151.001302, pixel 0.271800, grad 6.837004, laplacian 1441.615771, dIdt 17923.748633  
Epoch 6, Epoch SSIM: pixel 0.127570, grad 0.048240, laplacian 0.000004, dIdt 0.003150  
10it [00:25, 2.53s/it]

Epoch 7, Epoch loss: total 144.870363, pixel 0.272695, grad 7.092230, laplacian 1377.754041, dIdt 17922.755762  
Epoch 7, Epoch SSIM: pixel 0.114187, grad 0.049100, laplacian 0.000004, dIdt 0.003397  
10it [00:25, 2.55s/it]

Epoch 8, Epoch loss: total 144.143900, pixel 0.271950, grad 6.656942, laplacian 1374.842346, dIdt 17923.143945  
Epoch 8, Epoch SSIM: pixel 0.124894, grad 0.063830, laplacian 0.000004, dIdt 0.003758  
10it [00:25, 2.57s/it]

Epoch 9, Epoch loss: total 305.923668, pixel 0.250375, grad 6.159220, laplacian 2997.619458, dIdt 17924.524121  
Epoch 9, Epoch SSIM: pixel 0.160711, grad 0.081737, laplacian 0.000003, dIdt 0.003438  
10it [00:28, 2.81s/it]

Epoch 10, Epoch loss: total 229.865292, pixel 0.270687, grad 7.171700, laplacian 2226.908813, dIdt 17925.224414  
Epoch 10, Epoch SSIM: pixel 0.130257, grad 0.055432, laplacian 0.000003, dIdt 0.003692  
10it [00:26, 2.61s/it]

Epoch 11, Epoch loss: total 136.622093, pixel 0.272660, grad 6.719372, laplacian 1298.999915, dIdt 17924.213965  
Epoch 11, Epoch SSIM: pixel 0.113547, grad 0.067140, laplacian 0.000002, dIdt 0.003150  
10it [00:25, 2.55s/it]

Epoch 12, Epoch loss: total 382.101712, pixel 0.244614, grad 5.672535, laplacian 3764.267285, dIdt 17922.825879  
Epoch 12, Epoch SSIM: pixel 0.198117, grad 0.097671, laplacian 0.000002, dIdt 0.001798  
10it [00:25, 2.58s/it]

Epoch 13, Epoch loss: total 486.522043, pixel 0.250277, grad 6.514315, laplacian 4800.052246, dIdt 17922.826367  
Epoch 13, Epoch SSIM: pixel 0.178635, grad 0.072523, laplacian 0.000001, dIdt 0.001957  
10it [00:25, 2.52s/it]

Epoch 14, Epoch loss: total 202.811334, pixel 0.276339, grad 7.330801, laplacian 1954.777625, dIdt 17923.857422  
Epoch 14, Epoch SSIM: pixel 0.096058, grad 0.052143, laplacian 0.000002, dIdt 0.002490  
10it [00:25, 2.53s/it]

Epoch 15, Epoch loss: total 208.771985, pixel 0.265436, grad 6.119286, laplacian 2026.500403, dIdt 17921.657324  
Epoch 15, Epoch SSIM: pixel 0.131113, grad 0.083950, laplacian 0.000002, dIdt 0.002041  
10it [00:25, 2.58s/it]

Epoch 16, Epoch loss: total 672.771747, pixel 0.236812, grad 5.251298, laplacian 6675.180615, dIdt 17922.410742  
Epoch 16, Epoch SSIM: pixel 0.192263, grad 0.121163, laplacian 0.000001, dIdt 0.001410  
10it [00:25, 2.52s/it]

Epoch 17, Epoch loss: total 500.858636, pixel 0.254457, grad 6.946769, laplacian 4939.093286, dIdt 17919.806641  
Epoch 17, Epoch SSIM: pixel 0.132466, grad 0.073124, laplacian 0.000001, dIdt 0.001824  
10it [00:26, 2.62s/it]

Epoch 18, Epoch loss: total 224.621533, pixel 0.275744, grad 6.920702, laplacian 2176.980688, dIdt 17922.766504  
Epoch 18, Epoch SSIM: pixel 0.095047, grad 0.065678, laplacian 0.000002, dIdt 0.002811  
10it [00:25, 2.54s/it]

Epoch 19, Epoch loss: total 415.270639, pixel 0.259027, grad 5.761030, laplacian 4095.070068, dIdt 17920.496680  
Epoch 19, Epoch SSIM: pixel 0.142183, grad 0.102761, laplacian 0.000001, dIdt 0.001561  
10it [00:29, 2.90s/it]

Epoch 20, Epoch loss: total 873.464124, pixel 0.242809, grad 5.472133, laplacian 8679.895654, dIdt 17922.616504  
Epoch 20, Epoch SSIM: pixel 0.179072, grad 0.112573, laplacian 0.000001, dIdt 0.001012  
-----Finished-----  
-----Generating Data-----

100%| | 10/10 [00:00<00:00, 1203.84it/s]

-----Finished-----

10it [00:15, 1.58s/it]

ffmpeg version 4.1.9-0+deb10u1 Copyright (c) 2000-2022 the FFmpeg developers  
built with gcc 8 (Debian 8.3.0-6)

configuration: --prefix=/usr --extra-version=0+deb10u1 --toolchain=hardened  
--libdir=/usr/lib/x86\_64-linux-gnu --incdir=/usr/include/x86\_64-linux-gnu  
--arch=amd64 --enable-gpl --disable-stripping --enable-avresample --disable-  
filter=resample --enable-avisynth --enable-gnutls --enable-ladspa --enable-  
libaom --enable-libass --enable-libbluray --enable-libbs2b --enable-libcaca  
--enable-libcdio --enable-libcodec2 --enable-libflite --enable-libfontconfig  
--enable-libfreetype --enable-libfribidi --enable-libgme --enable-libgsm  
--enable-libjack --enable-libmp3lame --enable-libmysofa --enable-libopenjpeg  
--enable-libopenmpt --enable-libopus --enable-libpulse --enable-librsvg  
--enable-librubberband --enable-libshine --enable-libsnappy --enable-libsoxr  
--enable-libspeex --enable-libssh --enable-libtheora --enable-libtwolame  
--enable-libvidstab --enable-libvorbis --enable-libvpx --enable-libwavpack  
--enable-libwebp --enable-libx265 --enable-libxml2 --enable-libxvid --enable-  
libzmq --enable-libzvbi --enable-lv2 --enable-omx --enable-opengl --enable-  
opengl --enable-sdl2 --enable-libdc1394 --enable-libdrm --enable-libiec61883  
--enable-chromaprint --enable-frei0r --enable-libx264 --enable-shared

libavutil	56. 22.100 / 56. 22.100
libavcodec	58. 35.100 / 58. 35.100
libavformat	58. 20.100 / 58. 20.100
libavdevice	58. 5.100 / 58. 5.100
libavfilter	7. 40.101 / 7. 40.101
libavresample	4. 0. 0 / 4. 0. 0
libswscale	5. 3.100 / 5. 3.100
libswresample	3. 3.100 / 3. 3.100
libpostproc	55. 3.100 / 55. 3.100

Input #0, image2, from 'tmp/file%02d.png':

Duration: 00:00:05.00, start: 0.000000, bitrate: N/A

Stream #0:0: Video: png, rgba(pc), 1296x432 [SAR 2835:2835 DAR 3:1], 2 fps,  
2 tbr, 2 tbn, 2 tbc

Stream mapping:

Stream #0:0 -> #0:0 (png (native) -> h264 (libx264))

Press [q] to stop, [?] for help

[libx264 @ 0x5634a8239f00] using SAR=1/1

[libx264 @ 0x5634a8239f00] using cpu capabilities: MMX2 SSE2Fast SSSE3 SSE4.2  
AVX FMA3 BMI2 AVX2

[libx264 @ 0x5634a8239f00] profile High, level 3.1

[libx264 @ 0x5634a8239f00] 264 - core 155 r2917 0a84d98 - H.264/MPEG-4 AVC codec  
- Copyleft 2003-2018 - <http://www.videolan.org/x264.html> - options: cabac=1  
ref=3 deblock=1:0:0 analyse=0x3:0x113 me=hex subme=7 psy=1 psy\_rd=1.00:0.00  
mixed\_ref=1 me\_range=16 chroma\_me=1 trellis=1 8x8dct=1 cqm=0 deadzone=21,11  
fast\_pskip=1 chroma\_qp\_offset=-2 threads=3 lookahead\_threads=1 sliced\_threads=0  
nr=0 decimate=1 interlaced=0 bluray\_compat=0 constrained\_intra=0 bframes=3

```

b_pyramid=2 b_adapt=1 b_bias=0 direct=1 weightb=1 open_gop=0 weightp=2
keyint=250 keyint_min=25 scenecut=40 intra_refresh=0 rc_lookahead=40 rc=crf
mbtree=1 crf=23.0 qcomp=0.60 qpmin=0 qpmax=69 qpstep=4 ip_ratio=1.40 aq=1:1.00
Output #0, mp4, to '/home/jupyter/videos/runs/cameraman/experiments/0.01_1.0_0.1
_0.0/siren_uniformlr_1e-04_video.mp4':

```

Metadata:

encoder : Lavf58.20.100

Stream #0:0: Video: h264 (libx264) (avc1 / 0x31637661), yuv420p, 1296x432  
[SAR 1:1 DAR 3:1], q=-1--1, 30 fps, 15360 tbn, 30 tbc

Metadata:

encoder : Lavc58.35.100 libx264

Side data:

cpb: bitrate max/min/avg: 0/0/0 buffer size: 0 vbv\_delay: -1  
frame= 150 fps=103 q=-1.0 Lsize= 250kB time=00:00:04.90 bitrate=  
417.9kb/s dup=140 drop=0 speed=3.37x  
video:247kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB muxing  
overhead: 1.046852%

```

[libx264 @ 0x5634a8239f00] frame I:1      Avg QP:15.96  size: 50966
[libx264 @ 0x5634a8239f00] frame P:38    Avg QP:19.14  size:  5071
[libx264 @ 0x5634a8239f00] frame B:111   Avg QP:14.80  size:    81
[libx264 @ 0x5634a8239f00] consecutive B-frames:  1.3%  0.0%  0.0% 98.7%
[libx264 @ 0x5634a8239f00] mb I  I16..4: 52.2% 24.4% 23.5%
[libx264 @ 0x5634a8239f00] mb P  I16..4:  0.5%  1.5%  0.4%  P16..4:  4.2%  1.3%
1.3%  0.0%  0.0%   skip:90.7%
[libx264 @ 0x5634a8239f00] mb B  I16..4:  0.1%  0.0%  0.0%  B16..8:  2.7%  0.0%
0.0%  direct: 0.0%  skip:97.2%  L0:45.9% L1:54.1% BI: 0.1%
[libx264 @ 0x5634a8239f00] 8x8 transform intra:41.6% inter:65.7%
[libx264 @ 0x5634a8239f00] coded y,uvDC,uvAC intra: 50.5% 54.0% 50.6% inter:
1.0% 1.4% 0.8%
[libx264 @ 0x5634a8239f00] i16 v,h,dc,p: 69% 23%  8%  1%
[libx264 @ 0x5634a8239f00] i8 v,h,dc,ddl,ddr,vr,hd,vl,hu: 18% 11% 23%  6%  8%
9%  7%  7%  9%
[libx264 @ 0x5634a8239f00] i4 v,h,dc,ddl,ddr,vr,hd,vl,hu: 25% 24% 19%  5%  6%
6%  5%  5%  5%
[libx264 @ 0x5634a8239f00] i8c dc,h,v,p: 58% 19% 18%  6%
[libx264 @ 0x5634a8239f00] Weighted P-Frames: Y:0.0% UV:0.0%
[libx264 @ 0x5634a8239f00] ref P L0: 72.1% 21.6%  5.7%  0.6%
[libx264 @ 0x5634a8239f00] ref B L0: 72.7% 26.8%  0.5%
[libx264 @ 0x5634a8239f00] ref B L1: 97.6%  2.4%
[libx264 @ 0x5634a8239f00] kb/s:404.23

```

finished experiment # 2

-----Generating Data-----

100%| | 10/10 [00:00<00:00, 1250.20it/s]

-----Finished-----

-----Begin Training-----

10it [00:25, 2.52s/it]



Epoch 1, Epoch loss: total 26.621803, pixel 0.293029, grad 7.418119, laplacian 124.365458, dIdt 17930.726074  
Epoch 1, Epoch SSIM: pixel 0.144434, grad 0.034668, laplacian 0.000018, dIdt 0.001226  
10it [00:25, 2.58s/it]

Epoch 2, Epoch loss: total 32.793078, pixel 0.270055, grad 6.800129, laplacian 803.625409, dIdt 17929.688965  
Epoch 2, Epoch SSIM: pixel 0.171307, grad 0.047028, laplacian 0.000006, dIdt 0.002169  
10it [00:25, 2.52s/it]

Epoch 3, Epoch loss: total 34.653031, pixel 0.267791, grad 7.240684, laplacian 946.015164, dIdt 17925.415527  
Epoch 3, Epoch SSIM: pixel 0.132266, grad 0.037886, laplacian 0.000005, dIdt 0.002539  
10it [00:26, 2.61s/it]

Epoch 4, Epoch loss: total 32.987681, pixel 0.282658, grad 7.398408, laplacian 763.746381, dIdt 17923.541895  
Epoch 4, Epoch SSIM: pixel 0.074605, grad 0.036620, laplacian 0.000005, dIdt 0.002800  
10it [00:25, 2.54s/it]

Epoch 5, Epoch loss: total 36.060338, pixel 0.269357, grad 6.874941, laplacian 1123.810162, dIdt 17920.359180  
Epoch 5, Epoch SSIM: pixel 0.127593, grad 0.042870, laplacian 0.000003, dIdt 0.002469  
10it [00:25, 2.56s/it]

Epoch 6, Epoch loss: total 43.156741, pixel 0.265568, grad 6.978813, laplacian 1823.233777, dIdt 17919.032813  
Epoch 6, Epoch SSIM: pixel 0.131061, grad 0.048071, laplacian 0.000002, dIdt 0.003255  
10it [00:25, 2.54s/it]

Epoch 7, Epoch loss: total 35.063009, pixel 0.278535, grad 7.116309, laplacian 1000.033496, dIdt 17918.510645  
Epoch 7, Epoch SSIM: pixel 0.129885, grad 0.048056, laplacian 0.000004, dIdt 0.003179  
10it [00:25, 2.53s/it]

Epoch 8, Epoch loss: total 43.848952, pixel 0.261201, grad 6.329889, laplacian 1957.696680, dIdt 17915.975293  
Epoch 8, Epoch SSIM: pixel 0.171981, grad 0.073916, laplacian 0.000002, dIdt 0.002677  
10it [00:26, 2.62s/it]

Epoch 9, Epoch loss: total 53.777078, pixel 0.261862, grad 6.611146, laplacian 2922.459351, dIdt 17915.150879  
Epoch 9, Epoch SSIM: pixel 0.156788, grad 0.069717, laplacian 0.000002, dIdt 0.002660  
10it [00:27, 2.80s/it]

Epoch 10, Epoch loss: total 42.246780, pixel 0.274739, grad 7.115699, laplacian 1718.996826, dIdt 17913.637695  
Epoch 10, Epoch SSIM: pixel 0.087605, grad 0.057229, laplacian 0.000002, dIdt 0.003546  
10it [00:25, 2.57s/it]

Epoch 11, Epoch loss: total 47.746038, pixel 0.260828, grad 6.159050, laplacian 2364.864758, dIdt 17912.257617  
Epoch 11, Epoch SSIM: pixel 0.137628, grad 0.085086, laplacian 0.000002, dIdt 0.002645  
10it [00:25, 2.52s/it]

Epoch 12, Epoch loss: total 72.796843, pixel 0.245811, grad 5.922017, laplacian 4894.022437, dIdt 17910.020410  
Epoch 12, Epoch SSIM: pixel 0.164734, grad 0.095903, laplacian 0.000002, dIdt 0.002601  
10it [00:26, 2.62s/it]

Epoch 13, Epoch loss: total 50.957817, pixel 0.272831, grad 7.106015, laplacian 2591.329468, dIdt 17911.224707  
Epoch 13, Epoch SSIM: pixel 0.130631, grad 0.058398, laplacian 0.000002, dIdt 0.002785  
10it [00:25, 2.51s/it]

Epoch 14, Epoch loss: total 50.526639, pixel 0.266455, grad 6.305798, laplacian 2628.300171, dIdt 17911.192090  
Epoch 14, Epoch SSIM: pixel 0.155778, grad 0.079204, laplacian 0.000002, dIdt 0.002232  
10it [00:25, 2.54s/it]

Epoch 15, Epoch loss: total 76.133932, pixel 0.250388, grad 5.790076, laplacian 5241.235937, dIdt 17906.457910  
Epoch 15, Epoch SSIM: pixel 0.167398, grad 0.102950, laplacian 0.000001, dIdt 0.001513  
10it [00:25, 2.57s/it]

Epoch 16, Epoch loss: total 65.769697, pixel 0.263554, grad 6.825569, laplacian 4101.187622, dIdt 17905.896484  
Epoch 16, Epoch SSIM: pixel 0.152439, grad 0.064349, laplacian 0.000001, dIdt 0.002249  
10it [00:25, 2.51s/it]

```

Epoch 17, Epoch loss: total 50.029660, pixel 0.268612, grad 6.552466, laplacian
2554.009302, dIdt 17910.239844
Epoch 17, Epoch SSIM: pixel 0.137277, grad 0.075862, laplacian 0.000001, dIdt
0.002381

10it [00:25, 2.58s/it]

Epoch 18, Epoch loss: total 84.075417, pixel 0.241623, grad 5.531142, laplacian
6061.698853, dIdt 17903.125000
Epoch 18, Epoch SSIM: pixel 0.208097, grad 0.115840, laplacian 0.000001, dIdt
0.001382

10it [00:25, 2.54s/it]

Epoch 19, Epoch loss: total 90.826101, pixel 0.249573, grad 6.195786, laplacian
6670.383716, dIdt 17901.524512
Epoch 19, Epoch SSIM: pixel 0.192020, grad 0.089923, laplacian 0.000001, dIdt
0.001848

10it [00:28, 2.86s/it]

Epoch 20, Epoch loss: total 56.914589, pixel 0.272576, grad 6.785977, laplacian
3219.546289, dIdt 17905.890723
Epoch 20, Epoch SSIM: pixel 0.133891, grad 0.067046, laplacian 0.000001, dIdt
0.002562
-----Finished-----
-----Generating Data-----

100%|      | 10/10 [00:00<00:00, 1287.90it/s]

-----Finished-----

0it [00:00, ?it/s]/opt/conda/lib/python3.7/site-
packages/ipykernel_launcher.py:24: RuntimeWarning: More than 20 figures have
been opened. Figures created through the pyplot interface
(`matplotlib.pyplot.figure`) are retained until explicitly closed and may
consume too much memory. (To control this warning, see the rcParam
`figure.max_open_warning`).
10it [00:15, 1.58s/it]
ffmpeg version 4.1.9-0+deb10u1 Copyright (c) 2000-2022 the FFmpeg developers
built with gcc 8 (Debian 8.3.0-6)
configuration: --prefix=/usr --extra-version=0+deb10u1 --toolchain=hardened
--libdir=/usr/lib/x86_64-linux-gnu --incdir=/usr/include/x86_64-linux-gnu
--arch=amd64 --enable-gpl --disable-stripping --enable-avresample --disable-
filter=resample --enable-avisynth --enable-gnutls --enable-ladspa --enable-
libaom --enable-libass --enable-libbluray --enable-libbs2b --enable-libcaca
--enable-libcdio --enable-libcodec2 --enable-libflite --enable-libfontconfig
--enable-libfreetype --enable-libfribidi --enable-libgme --enable-libgsm
--enable-libjack --enable-libmp3lame --enable-libmysofa --enable-libopenjpeg
--enable-libopenmpt --enable-libopus --enable-libpulse --enable-librsvg
--enable-librubberband --enable-libshine --enable-libsnappy --enable-libsoxr
--enable-libspeex --enable-libssh --enable-libtheora --enable-libtwolame

```

```

--enable-libvidstab --enable-libvorbis --enable-libvpx --enable-libwavpack
--enable-libwebp --enable-libx265 --enable-libxml2 --enable-libxvid --enable-
libzmq --enable-libzvbi --enable-lv2 --enable-omx --enable-openal --enable-
opengl --enable-sdl2 --enable-libdc1394 --enable-libdrm --enable-libiec61883
--enable-chromaprint --enable-frei0r --enable-libx264 --enable-shared
  libavutil      56. 22.100 / 56. 22.100
  libavcodec      58. 35.100 / 58. 35.100
  libavformat      58. 20.100 / 58. 20.100
  libavdevice      58.  5.100 / 58.  5.100
  libavfilter       7. 40.101 /  7. 40.101
  libavresample     4.  0.  0 /  4.  0.  0
  libswscale        5.  3.100 /  5.  3.100
  libswresample     3.  3.100 /  3.  3.100
  libpostproc      55.  3.100 / 55.  3.100
Input #0, image2, from 'tmp/file%02d.png':
  Duration: 00:00:05.00, start: 0.000000, bitrate: N/A
    Stream #0:0: Video: png, rgba(pc), 1296x432 [SAR 2835:2835 DAR 3:1], 2 fps,
2 tbr, 2 tbn, 2 tbc
Stream mapping:
  Stream #0:0 -> #0:0 (png (native) -> h264 (libx264))
Press [q] to stop, [?] for help
[libx264 @ 0x561f8d9d5f00] using SAR=1/1
[libx264 @ 0x561f8d9d5f00] using cpu capabilities: MMX2 SSE2Fast SSSE3 SSE4.2
AVX FMA3 BMI2 AVX2
[libx264 @ 0x561f8d9d5f00] profile High, level 3.1
[libx264 @ 0x561f8d9d5f00] 264 - core 155 r2917 0a84d98 - H.264/MPEG-4 AVC codec
- Copyleft 2003-2018 - http://www.videolan.org/x264.html - options: cabac=1
ref=3 deblock=1:0:0 analyse=0x3:0x113 me=hex subme=7 psy=1 psy_rd=1.00:0.00
mixed_ref=1 me_range=16 chroma_me=1 trellis=1 8x8dct=1 cqm=0 deadzone=21,11
fast_pskip=1 chroma_qp_offset=-2 threads=3 lookahead_threads=1 sliced_threads=0
nr=0 decimate=1 interlaced=0 bluray_compat=0 constrained_intra=0 bframes=3
b_pyramid=2 b_adapt=1 b_bias=0 direct=1 weightb=1 open_gop=0 weightp=2
keyint=250 keyint_min=25 scenecut=40 intra_refresh=0 rc_lookahead=40 rc=crf
mbtree=1 crf=23.0 qcomp=0.60 qpmin=0 qpmax=69 qpstep=4 ip_ratio=1.40 aq=1:1.00
Output #0, mp4, to '/home/jupyter/videos/runs/cameraman/experiments/0.1_1.0_0.01
_0.001/siren_uniformlr_1e-04_video.mp4':
  Metadata:
    encoder      : Lavf58.20.100
    Stream #0:0: Video: h264 (libx264) (avc1 / 0x31637661), yuv420p, 1296x432
[SAR 1:1 DAR 3:1], q=-1--1, 30 fps, 15360 tbn, 30 tbc
  Metadata:
    encoder      : Lavc58.35.100 libx264
  Side data:
    cpb: bitrate max/min/avg: 0/0/0 buffer size: 0 vbv_delay: -1
frame= 150 fps=104 q=-1.0 Lsize=      260kB time=00:00:04.90 bitrate=
434.8kb/s dup=140 drop=0 speed=3.39x
video:258kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB muxing
overhead: 1.005734%

```

```

[libx264 @ 0x561f8d9d5f00] frame I:1      Avg QP:16.82  size: 51330
[libx264 @ 0x561f8d9d5f00] frame P:38    Avg QP:19.06  size:  5331
[libx264 @ 0x561f8d9d5f00] frame B:111   Avg QP:14.84  size:    82
[libx264 @ 0x561f8d9d5f00] consecutive B-frames:  1.3%  0.0%  0.0% 98.7%
[libx264 @ 0x561f8d9d5f00] mb I   I16..4: 50.7% 26.2% 23.1%
[libx264 @ 0x561f8d9d5f00] mb P   I16..4:  0.5%  1.6%  0.4%  P16..4:  4.2%  1.4%
1.3%  0.0%  0.0%    skip:90.7%
[libx264 @ 0x561f8d9d5f00] mb B   I16..4:  0.1%  0.0%  0.0%  B16..8:  2.8%  0.0%
0.0% direct: 0.0% skip:97.1% L0:47.8% L1:52.1% BI:  0.1%
[libx264 @ 0x561f8d9d5f00] 8x8 transform intra:44.0% inter:67.3%
[libx264 @ 0x561f8d9d5f00] coded y,uvDC,uvAC intra: 52.8% 54.8% 52.0% inter:
1.0% 1.4% 0.8%
[libx264 @ 0x561f8d9d5f00] i16 v,h,dc,p: 69% 22%  8%  0%
[libx264 @ 0x561f8d9d5f00] i8  v,h,dc,ddl,ddr,vr,hd,vl,hu: 19% 11% 21%  7%  8%
9%  7%  8% 10%
[libx264 @ 0x561f8d9d5f00] i4  v,h,dc,ddl,ddr,vr,hd,vl,hu: 25% 24% 18%  6%  5%
6%  5%  5%  5%
[libx264 @ 0x561f8d9d5f00] i8c dc,h,v,p: 58% 18% 18%  6%
[libx264 @ 0x561f8d9d5f00] Weighted P-Frames: Y:0.0% UV:0.0%
[libx264 @ 0x561f8d9d5f00] ref P L0: 70.8% 22.2%  6.3%  0.7%
[libx264 @ 0x561f8d9d5f00] ref B L0: 68.6% 31.1%  0.4%
[libx264 @ 0x561f8d9d5f00] ref B L1: 97.5%  2.5%
[libx264 @ 0x561f8d9d5f00] kb/s:420.80

```

finished experiment # 3

-----Generating Data-----

100%| | 10/10 [00:00<00:00, 1108.17it/s]

-----Finished-----

-----Begin Training-----

10it [00:24, 2.49s/it]

Epoch 1, Epoch loss: total 1794.866235, pixel 0.344890, grad 8.248509, laplacian 109.372216, dIdt 17947.564844

Epoch 1, Epoch SSIM: pixel 0.101625, grad 0.026833, laplacian 0.000026, dIdt 0.001582

10it [00:25, 2.60s/it]

Epoch 2, Epoch loss: total 1775.744592, pixel 0.408818, grad 11.107056, laplacian 1666.020834, dIdt 17740.780859

Epoch 2, Epoch SSIM: pixel 0.007082, grad 0.014969, laplacian 0.000005, dIdt 0.000191

10it [00:25, 2.51s/it]

Epoch 3, Epoch loss: total 1809.437537, pixel 0.349068, grad 14.187024, laplacian 10766.420215, dIdt 17986.707422

Epoch 3, Epoch SSIM: pixel 0.042241, grad 0.014762, laplacian 0.000002, dIdt 0.000714

10it [00:26, 2.61s/it]

Epoch 4, Epoch loss: total 2100.757666, pixel 0.343609, grad 30.208072,  
laplacian 304889.844873, dIdt 17958.674316

Epoch 4, Epoch SSIM: pixel 0.040865, grad 0.013085, laplacian 0.000003, dIdt  
0.000097

10it [00:25, 2.54s/it]

Epoch 5, Epoch loss: total 2041.692188, pixel 0.329382, grad 25.694508,  
laplacian 249898.331641, dIdt 17917.934668

Epoch 5, Epoch SSIM: pixel 0.050619, grad 0.010567, laplacian 0.000001, dIdt  
-0.000011

10it [00:25, 2.53s/it]

Epoch 6, Epoch loss: total 1988.387708, pixel 0.346298, grad 26.383153,  
laplacian 195183.581250, dIdt 17932.037793

Epoch 6, Epoch SSIM: pixel 0.041036, grad 0.009976, laplacian 0.000001, dIdt  
-0.000056

10it [00:25, 2.56s/it]

Epoch 7, Epoch loss: total 2053.435474, pixel 0.348727, grad 28.968291,  
laplacian 263050.827734, dIdt 17903.842773

Epoch 7, Epoch SSIM: pixel 0.061297, grad 0.008113, laplacian 0.000000, dIdt  
-0.000084

10it [00:25, 2.52s/it]

Epoch 8, Epoch loss: total 2202.734277, pixel 0.350646, grad 37.785837,  
laplacian 414788.740625, dIdt 17879.451172

Epoch 8, Epoch SSIM: pixel 0.066285, grad 0.009018, laplacian 0.000000, dIdt  
-0.000071

10it [00:26, 2.61s/it]

Epoch 9, Epoch loss: total 2238.020190, pixel 0.368077, grad 42.810662,  
laplacian 452185.175781, dIdt 17858.345410

Epoch 9, Epoch SSIM: pixel 0.065114, grad 0.007509, laplacian 0.000000, dIdt  
-0.000064

10it [00:28, 2.80s/it]

Epoch 10, Epoch loss: total 2253.463062, pixel 0.402176, grad 44.785559,  
laplacian 471901.925000, dIdt 17815.606543

Epoch 10, Epoch SSIM: pixel 0.080571, grad 0.006661, laplacian 0.000000, dIdt  
-0.000145

10it [00:25, 2.56s/it]

Epoch 11, Epoch loss: total 2131.985803, pixel 0.458949, grad 46.644411,  
laplacian 349151.180469, dIdt 17828.341211

Epoch 11, Epoch SSIM: pixel 0.079338, grad 0.007240, laplacian 0.000000, dIdt  
-0.000152

10it [00:25, 2.52s/it]

Epoch 12, Epoch loss: total 2197.925232, pixel 0.499348, grad 45.420939,  
laplacian 419567.213281, dIdt 17783.574902

Epoch 12, Epoch SSIM: pixel 0.109144, grad 0.007339, laplacian 0.000000, dIdt  
-0.000089

10it [00:25, 2.52s/it]

Epoch 13, Epoch loss: total 2217.877588, pixel 0.552990, grad 55.927468,  
laplacian 430840.253125, dIdt 17870.367188

Epoch 13, Epoch SSIM: pixel 0.068456, grad 0.006207, laplacian 0.000000, dIdt  
-0.000125

10it [00:25, 2.57s/it]

Epoch 14, Epoch loss: total 2149.671509, pixel 0.526460, grad 45.183729,  
laplacian 371920.375781, dIdt 17777.505762

Epoch 14, Epoch SSIM: pixel 0.091550, grad 0.006252, laplacian 0.000000, dIdt  
-0.000080

10it [00:25, 2.51s/it]

Epoch 15, Epoch loss: total 2254.380457, pixel 0.643062, grad 57.125675,  
laplacian 476514.600000, dIdt 17778.651660

Epoch 15, Epoch SSIM: pixel 0.072173, grad 0.005382, laplacian 0.000000, dIdt  
-0.000172

10it [00:25, 2.60s/it]

Epoch 16, Epoch loss: total 2263.104126, pixel 0.638717, grad 59.793797,  
laplacian 489054.079688, dIdt 17740.493164

Epoch 16, Epoch SSIM: pixel 0.087960, grad 0.006100, laplacian 0.000000, dIdt  
-0.000144

10it [00:25, 2.51s/it]

Epoch 17, Epoch loss: total 2363.737866, pixel 0.753830, grad 69.708949,  
laplacian 596823.770312, dIdt 17669.132910

Epoch 17, Epoch SSIM: pixel 0.079889, grad 0.005548, laplacian 0.000000, dIdt  
-0.000092

10it [00:26, 2.61s/it]

Epoch 18, Epoch loss: total 2445.548071, pixel 0.870792, grad 76.627430,  
laplacian 681998.400000, dIdt 17635.487500

Epoch 18, Epoch SSIM: pixel 0.079763, grad 0.004916, laplacian -0.000000, dIdt  
-0.000127

10it [00:25, 2.54s/it]

Epoch 19, Epoch loss: total 2680.395239, pixel 0.880309, grad 105.306596,  
laplacian 926203.262500, dIdt 17541.909961

Epoch 19, Epoch SSIM: pixel 0.066640, grad 0.004182, laplacian 0.000000, dIdt  
-0.000085

```

10it [00:27, 2.80s/it]

Epoch 20, Epoch loss: total 2818.872852, pixel 1.090704, grad 107.666595,
laplacian 1066049.231250, dIdt 17528.225000
Epoch 20, Epoch SSIM: pixel 0.063074, grad 0.001992, laplacian -0.000000, dIdt
-0.000045
-----Finished-----
-----Generating Data-----

100%|      | 10/10 [00:00<00:00, 1285.45it/s]

-----Finished-----

10it [00:16, 1.61s/it]
ffmpeg version 4.1.9-0+deb10u1 Copyright (c) 2000-2022 the FFmpeg developers
  built with gcc 8 (Debian 8.3.0-6)
  configuration: --prefix=/usr --extra-version=0+deb10u1 --toolchain=hardened
--libdir=/usr/lib/x86_64-linux-gnu --incdir=/usr/include/x86_64-linux-gnu
--arch=amd64 --enable-gpl --disable-stripping --enable-avresample --disable-
filter=resample --enable-avisynth --enable-gnutls --enable-ladspa --enable-
libaom --enable-libass --enable-libbluray --enable-libbs2b --enable-libcaca
--enable-libcdio --enable-libcodec2 --enable-libflite --enable-libfontconfig
--enable-libfreetype --enable-libfribidi --enable-libgme --enable-libgsm
--enable-libjack --enable-libmp3lame --enable-libmysofa --enable-libopenjpeg
--enable-libopenmpt --enable-libopus --enable-libpulse --enable-librsvg
--enable-librubberband --enable-libshine --enable-libsnappy --enable-libsoxr
--enable-libspeex --enable-libssh --enable-libtheora --enable-libtwolame
--enable-libvidstab --enable-libvorbis --enable-libvpx --enable-libwavpack
--enable-libwebp --enable-libx265 --enable-libxml2 --enable-libxvid --enable-
libzmq --enable-libzvbi --enable-lv2 --enable-omx --enable-openal --enable-
opengl --enable-sdl2 --enable-libdc1394 --enable-libdrm --enable-libiec61883
--enable-chromaprint --enable-frei0r --enable-libx264 --enable-shared
   libavutil      56. 22.100 / 56. 22.100
   libavcodec     58. 35.100 / 58. 35.100
   libavformat    58. 20.100 / 58. 20.100
   libavdevice    58.  5.100 / 58.  5.100
   libavfilter     7. 40.101 /  7. 40.101
   libavresample   4.  0.  0 /  4.  0.  0
   libswscale      5.  3.100 /  5.  3.100
   libswresample   3.  3.100 /  3.  3.100
   libpostproc    55.  3.100 / 55.  3.100
Input #0, image2, from 'tmp/file%02d.png':
  Duration: 00:00:05.00, start: 0.000000, bitrate: N/A
    Stream #0:0: Video: png, rgba(pc), 1296x432 [SAR 2835:2835 DAR 3:1], 2 fps,
2 tbr, 2 tbn, 2 tbc
Stream mapping:
  Stream #0:0 -> #0:0 (png (native) -> h264 (libx264))
Press [q] to stop, [?] for help
[libx264 @ 0x555fe683af00] using SAR=1/1
[libx264 @ 0x555fe683af00] using cpu capabilities: MMX2 SSE2Fast SSSE3 SSE4.2

```



```

AVX FMA3 BMI2 AVX2
[libx264 @ 0x555fe683af00] profile High, level 3.1
[libx264 @ 0x555fe683af00] 264 - core 155 r2917 0a84d98 - H.264/MPEG-4 AVC codec
- Copyleft 2003-2018 - http://www.videolan.org/x264.html - options: cabac=1
ref=3 deblock=1:0:0 analyse=0x3:0x113 me=hex subme=7 psy=1 psy_rd=1.00:0.00
mixed_ref=1 me_range=16 chroma_me=1 trellis=1 8x8dct=1 cqm=0 deadzone=21,11
fast_pskip=1 chroma_qp_offset=-2 threads=3 lookahead_threads=1 sliced_threads=0
nr=0 decimate=1 interlaced=0 bluray_compat=0 constrained_intra=0 bframes=3
b_pyramid=2 b_adapt=1 b_bias=0 direct=1 weightb=1 open_gop=0 weightp=2
keyint=250 keyint_min=25 scenecut=40 intra_refresh=0 rc_lookahead=40 rc=crf
mbtree=1 crf=23.0 qcomp=0.60 qpmin=0 qpmax=69 qpstep=4 ip_ratio=1.40 aq=1:1.00
Output #0, mp4, to '/home/jupyter/videos/runs/cameraman/experiments/0.001_0.0_0.
001_0.1/siren_uniformlr_1e-04_video.mp4':
  Metadata:
    encoder           : Lavf58.20.100
    Stream #0:0: Video: h264 (libx264) (avc1 / 0x31637661), yuv420p, 1296x432
[SAR 1:1 DAR 3:1], q=-1--1, 30 fps, 15360 tbn, 30 tbc
    Metadata:
      encoder          : Lavc58.35.100 libx264
    Side data:
      cpb: bitrate max/min/avg: 0/0/0 buffer size: 0 vbv_delay: -1
frame= 150 fps=103 q=-1.0 Lsize=      267kB time=00:00:04.90 bitrate=
445.9kb/s dup=140 drop=0 speed=3.36x
video:264kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB muxing
overhead: 0.980566%
[libx264 @ 0x555fe683af00] frame I:1      Avg QP:16.81  size: 51178
[libx264 @ 0x555fe683af00] frame P:38     Avg QP:19.13  size:  5514
[libx264 @ 0x555fe683af00] frame B:111    Avg QP:14.75  size:    82
[libx264 @ 0x555fe683af00] consecutive B-frames:  1.3%  0.0%  0.0% 98.7%
[libx264 @ 0x555fe683af00] mb I  I16..4: 51.1% 23.9% 25.0%
[libx264 @ 0x555fe683af00] mb P  I16..4:  0.6%  1.9%  1.1% P16..4:  3.8%  0.9%
1.0%  0.0%  0.0%  skip:90.6%
[libx264 @ 0x555fe683af00] mb B  I16..4:  0.1%  0.0%  0.0% B16..8:  2.8%  0.0%
0.0% direct: 0.0% skip:97.1% L0:46.7% L1:53.2% BI: 0.1%
[libx264 @ 0x555fe683af00] 8x8 transform intra:40.2% inter:59.5%
[libx264 @ 0x555fe683af00] coded y,uvDC,uvAC intra: 57.9% 62.8% 59.3% inter:
0.8% 1.1% 0.6%
[libx264 @ 0x555fe683af00] i16 v,h,dc,p: 70% 21%  8%  1%
[libx264 @ 0x555fe683af00] i8 v,h,dc,ddl,ddr,vr,hd,vl,hu: 18% 13% 19%  5% 10%
9% 10%  7%  9%
[libx264 @ 0x555fe683af00] i4 v,h,dc,ddl,ddr,vr,hd,vl,hu: 21% 22% 14%  5%  9%
8%  9%  5%  6%
[libx264 @ 0x555fe683af00] i8c dc,h,v,p: 57% 19% 17%  7%
[libx264 @ 0x555fe683af00] Weighted P-Frames: Y:0.0% UV:0.0%
[libx264 @ 0x555fe683af00] ref P L0: 73.7% 19.3%  6.4%  0.6%
[libx264 @ 0x555fe683af00] ref B L0: 70.0% 29.5%  0.4%
[libx264 @ 0x555fe683af00] ref B L1: 97.6%  2.4%
[libx264 @ 0x555fe683af00] kb/s:431.63

```

```

finished experiment # 4
-----Generating Data-----

100%|      | 10/10 [00:00<00:00, 1280.47it/s]

-----Finished-----
-----Begin Training-----

10it [00:24,  2.47s/it]

Epoch 1, Epoch loss: total 1814.866809, pixel 0.346538, grad 8.280921, laplacian
145.447729, dIdt 17916.945215
Epoch 1, Epoch SSIM: pixel 0.066716, grad 0.027082, laplacian 0.000023, dIdt
0.001077

10it [00:25,  2.58s/it]

Epoch 2, Epoch loss: total 1860.611877, pixel 0.380354, grad 9.547269, laplacian
526.857953, dIdt 17979.984668
Epoch 2, Epoch SSIM: pixel 0.026077, grad 0.020811, laplacian 0.000011, dIdt
0.000759

10it [00:25,  2.52s/it]

Epoch 3, Epoch loss: total 2007.487341, pixel 0.338607, grad 9.528469, laplacian
2000.423846, dIdt 17975.778418
Epoch 3, Epoch SSIM: pixel 0.052165, grad 0.020526, laplacian 0.000005, dIdt
0.000811

10it [00:25,  2.51s/it]

Epoch 4, Epoch loss: total 1915.231946, pixel 0.323519, grad 8.647112, laplacian
1095.440375, dIdt 17967.172852
Epoch 4, Epoch SSIM: pixel 0.066524, grad 0.026051, laplacian 0.000006, dIdt
0.001681

10it [00:25,  2.58s/it]

Epoch 5, Epoch loss: total 2000.744348, pixel 0.311763, grad 8.571822, laplacian
1984.589478, dIdt 17934.017676
Epoch 5, Epoch SSIM: pixel 0.034062, grad 0.021815, laplacian 0.000003, dIdt
0.001552

10it [00:25,  2.52s/it]

Epoch 6, Epoch loss: total 2007.131433, pixel 0.316525, grad 8.451786, laplacian
2063.998999, dIdt 17919.631836
Epoch 6, Epoch SSIM: pixel -0.000337, grad 0.021222, laplacian 0.000005, dIdt
0.001306

10it [00:26,  2.60s/it]

Epoch 7, Epoch loss: total 2127.891479, pixel 0.316967, grad 8.532283, laplacian
3310.638397, dIdt 17879.783398
Epoch 7, Epoch SSIM: pixel 0.048897, grad 0.022880, laplacian 0.000003, dIdt
0.001553

```

10it [00:25, 2.52s/it]

Epoch 8, Epoch loss: total 2045.645618, pixel 0.325292, grad 8.582417, laplacian 2518.142517, dIdt 17849.236426

Epoch 8, Epoch SSIM: pixel 0.111544, grad 0.018265, laplacian 0.000004, dIdt 0.001153

10it [00:26, 2.60s/it]

Epoch 9, Epoch loss: total 2195.640747, pixel 0.326631, grad 9.105421, laplacian 4041.611841, dIdt 17820.474316

Epoch 9, Epoch SSIM: pixel 0.149717, grad 0.019854, laplacian 0.000001, dIdt 0.000889

10it [00:27, 2.80s/it]

Epoch 10, Epoch loss: total 2338.577417, pixel 0.340185, grad 9.802977, laplacian 5436.615308, dIdt 17847.726270

Epoch 10, Epoch SSIM: pixel 0.144651, grad 0.016809, laplacian 0.000001, dIdt 0.000337

10it [00:25, 2.51s/it]

Epoch 11, Epoch loss: total 2447.192749, pixel 0.341265, grad 10.091131, laplacian 6479.799939, dIdt 17887.802930

Epoch 11, Epoch SSIM: pixel 0.174936, grad 0.015297, laplacian 0.000003, dIdt 0.000655

10it [00:25, 2.59s/it]

Epoch 12, Epoch loss: total 2319.233435, pixel 0.359956, grad 9.305313, laplacian 5283.228564, dIdt 17812.452246

Epoch 12, Epoch SSIM: pixel 0.179723, grad 0.019449, laplacian 0.000001, dIdt 0.000341

10it [00:25, 2.51s/it]

Epoch 13, Epoch loss: total 2448.509143, pixel 0.396323, grad 9.817775, laplacian 6651.720801, dIdt 17731.228809

Epoch 13, Epoch SSIM: pixel 0.197913, grad 0.019172, laplacian 0.000001, dIdt 0.000515

10it [00:25, 2.58s/it]

Epoch 14, Epoch loss: total 2672.038867, pixel 0.403108, grad 11.305337, laplacian 8852.258960, dIdt 17751.044727

Epoch 14, Epoch SSIM: pixel 0.158298, grad 0.013831, laplacian 0.000001, dIdt 0.000052

10it [00:25, 2.51s/it]

Epoch 15, Epoch loss: total 3214.529370, pixel 0.459044, grad 12.215571, laplacian 14272.570410, dIdt 17745.975977

Epoch 15, Epoch SSIM: pixel 0.171228, grad 0.014736, laplacian 0.000000, dIdt 0.000307

```

10it [00:25, 2.59s/it]

Epoch 16, Epoch loss: total 3742.840369, pixel 0.519699, grad 12.350855,
laplacian 19620.176904, dIdt 17679.520605
Epoch 16, Epoch SSIM: pixel 0.158672, grad 0.015579, laplacian 0.000000, dIdt
0.000220

10it [00:25, 2.52s/it]

Epoch 17, Epoch loss: total 3773.448828, pixel 0.454644, grad 13.829960,
laplacian 19860.530664, dIdt 17731.110742
Epoch 17, Epoch SSIM: pixel 0.179908, grad 0.015471, laplacian 0.000000, dIdt
0.000116

10it [00:25, 2.51s/it]

Epoch 18, Epoch loss: total 3802.279419, pixel 0.490274, grad 13.280104,
laplacian 20207.345020, dIdt 17677.744922
Epoch 18, Epoch SSIM: pixel 0.168784, grad 0.011985, laplacian 0.000000, dIdt
0.000094

10it [00:25, 2.58s/it]

Epoch 19, Epoch loss: total 3800.307642, pixel 0.537183, grad 13.228332,
laplacian 20240.468457, dIdt 17624.952148
Epoch 19, Epoch SSIM: pixel 0.174455, grad 0.014050, laplacian 0.000000, dIdt
0.000132

10it [00:27, 2.80s/it]

Epoch 20, Epoch loss: total 4174.798901, pixel 0.581214, grad 14.603018,
laplacian 24003.438037, dIdt 17592.708789
Epoch 20, Epoch SSIM: pixel 0.150035, grad 0.015564, laplacian 0.000000, dIdt
0.000074
-----Finished-----
-----Generating Data-----

100%|      | 10/10 [00:00<00:00, 1286.28it/s]

-----Finished-----

10it [00:21, 2.10s/it]
ffmpeg version 4.1.9-0+deb10u1 Copyright (c) 2000-2022 the FFmpeg developers
built with gcc 8 (Debian 8.3.0-6)
configuration: --prefix=/usr --extra-version=0+deb10u1 --toolchain=hardened
--libdir=/usr/lib/x86_64-linux-gnu --incdir=/usr/include/x86_64-linux-gnu
--arch=amd64 --enable-gpl --disable-stripping --enable-avresample --disable-
filter=resample --enable-avisynth --enable-gnutls --enable-ladspa --enable-
libaom --enable-libass --enable-libbluray --enable-libbs2b --enable-libcaca
--enable-libcdio --enable-libcodec2 --enable-libflite --enable-libfontconfig
--enable-libfreetype --enable-libfribidi --enable-libgme --enable-libgsm
--enable-libjack --enable-libmp3lame --enable-libmysofa --enable-libopenjpeg
--enable-libopenmpt --enable-libopus --enable-libpulse --enable-librsvg
--enable-librubberband --enable-libshine --enable-libsnappy --enable-libsoxr

```

```

--enable-libspeex --enable-libssh --enable-libtheora --enable-libtwolame
--enable-libvidstab --enable-libvorbis --enable-libvpx --enable-libwavpack
--enable-libwebp --enable-libx265 --enable-libxml2 --enable-libxvid --enable-
libzmq --enable-libzvbi --enable-lv2 --enable-omx --enable-openal --enable-
opengl --enable-sdl2 --enable-libdc1394 --enable-libdrm --enable-libiec61883
--enable-chromaprint --enable-frei0r --enable-libx264 --enable-shared
  libavutil      56. 22.100 / 56. 22.100
  libavcodec      58. 35.100 / 58. 35.100
  libavformat     58. 20.100 / 58. 20.100
  libavdevice     58.  5.100 / 58.  5.100
  libavfilter     7. 40.101 /  7. 40.101
  libavresample   4.  0.  0 /  4.  0.  0
  libswscale      5.  3.100 /  5.  3.100
  libswresample   3.  3.100 /  3.  3.100
  libpostproc    55.  3.100 / 55.  3.100
Input #0, image2, from 'tmp/file%02d.png':
  Duration: 00:00:05.00, start: 0.000000, bitrate: N/A
    Stream #0:0: Video: png, rgba(pc), 1296x432 [SAR 2835:2835 DAR 3:1], 2 fps,
2 tbr, 2 tbn, 2 tbc
Stream mapping:
  Stream #0:0 -> #0:0 (png (native) -> h264 (libx264))
Press [q] to stop, [?] for help
[libx264 @ 0x55a062731f00] using SAR=1/1
[libx264 @ 0x55a062731f00] using cpu capabilities: MMX2 SSE2Fast SSSE3 SSE4.2
AVX FMA3 BMI2 AVX2
[libx264 @ 0x55a062731f00] profile High, level 3.1
[libx264 @ 0x55a062731f00] 264 - core 155 r2917 0a84d98 - H.264/MPEG-4 AVC codec
- Copyleft 2003-2018 - http://www.videolan.org/x264.html - options: cabac=1
ref=3 deblock=1:0:0 analyse=0x3:0x113 me=hex subme=7 psy=1 psy_rd=1.00:0.00
mixed_ref=1 me_range=16 chroma_me=1 trellis=1 8x8dct=1 cqm=0 deadzone=21,11
fast_pskip=1 chroma_qp_offset=-2 threads=3 lookahead_threads=1 sliced_threads=0
nr=0 decimate=1 interlaced=0 bluray_compat=0 constrained_intra=0 bframes=3
b_pyramid=2 b_adapt=1 b_bias=0 direct=1 weightb=1 open_gop=0 weightp=2
keyint=250 keyint_min=25 scenecut=40 intra_refresh=0 rc_lookahead=40 rc=crf
mbtree=1 crf=23.0 qcomp=0.60 qpmin=0 qpmax=69 qpstep=4 ip_ratio=1.40 aq=1:1.00
Output #0, mp4, to '/home/jupyter/videos/runs/cameraman/experiments/1.0_1.0_0.1_
0.1/siren_uniformlr_1e-04_video.mp4':
  Metadata:
    encoder      : Lavf58.20.100
    Stream #0:0: Video: h264 (libx264) (avc1 / 0x31637661), yuv420p, 1296x432
[SAR 1:1 DAR 3:1], q=-1--1, 30 fps, 15360 tbn, 30 tbc
  Metadata:
    encoder      : Lavc58.35.100 libx264
  Side data:
    cpb: bitrate max/min/avg: 0/0/0 buffer size: 0 vbv_delay: -1
frame= 150 fps=104 q=-1.0 Lsize=      254kB time=00:00:04.90 bitrate=
424.9kbts/s dup=140 drop=0 speed= 3.4x
video:252kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB muxing

```

```

overhead: 1.029547%
[libx264 @ 0x55a062731f00] frame I:1      Avg QP:15.70  size: 47904
[libx264 @ 0x55a062731f00] frame P:38    Avg QP:19.00  size:  5259
[libx264 @ 0x55a062731f00] frame B:111   Avg QP:14.73  size:    82
[libx264 @ 0x55a062731f00] consecutive B-frames:  1.3%  0.0%  0.0% 98.7%
[libx264 @ 0x55a062731f00] mb I  I16..4: 58.3% 17.0% 24.7%
[libx264 @ 0x55a062731f00] mb P  I16..4:  0.6%  2.3%  0.9%  P16..4:  3.8%  0.9%
0.9%  0.0%  0.0%   skip:90.6%
[libx264 @ 0x55a062731f00] mb B  I16..4:  0.1%  0.0%  0.0%  B16..8:  2.9%  0.0%
0.0%  direct: 0.0%  skip:97.0%  L0:46.6% L1:53.3% BI: 0.1%
[libx264 @ 0x55a062731f00] 8x8 transform intra:42.0% inter:59.6%
[libx264 @ 0x55a062731f00] coded y,uvDC,uvAC intra: 58.4% 63.5% 60.5% inter:
0.8% 1.1% 0.6%
[libx264 @ 0x55a062731f00] i16 v,h,dc,p: 73% 20%  6%  1%
[libx264 @ 0x55a062731f00] i8 v,h,dc,ddl,ddr,vr,hd,vl,hu: 14% 13% 17%  7% 10%
10%  9%  9% 12%
[libx264 @ 0x55a062731f00] i4 v,h,dc,ddl,ddr,vr,hd,vl,hu: 25% 22% 16%  5%  7%
6%  7%  5%  6%
[libx264 @ 0x55a062731f00] i8c dc,h,v,p: 54% 20% 19%  8%
[libx264 @ 0x55a062731f00] Weighted P-Frames: Y:0.0% UV:0.0%
[libx264 @ 0x55a062731f00] ref P L0: 73.8% 18.4%  7.1%  0.7%
[libx264 @ 0x55a062731f00] ref B L0: 71.1% 28.5%  0.4%
[libx264 @ 0x55a062731f00] ref B L1: 97.3%  2.7%
[libx264 @ 0x55a062731f00] kb/s:411.04

finished experiment # 5
-----Generating Data-----

100%|      | 10/10 [00:00<00:00, 1159.22it/s]

-----Finished-----
-----Begin Training-----

10it [00:24,  2.48s/it]

Epoch 1, Epoch loss: total 190.568272, pixel 0.344656, grad 8.148641, laplacian
104.195838, dIdt 17933.348535
Epoch 1, Epoch SSIM: pixel 0.032766, grad 0.027506, laplacian 0.000024, dIdt
0.001690

10it [00:25,  2.60s/it]

Epoch 2, Epoch loss: total 223.934142, pixel 0.345734, grad 9.043941, laplacian
444.560037, dIdt 17857.340234
Epoch 2, Epoch SSIM: pixel 0.018716, grad 0.018599, laplacian 0.000010, dIdt
0.000977

10it [00:25,  2.51s/it]

Epoch 3, Epoch loss: total 255.237726, pixel 0.334775, grad 8.909799, laplacian
747.814856, dIdt 17956.492871

```

Epoch 3, Epoch SSIM: pixel 0.020453, grad 0.020034, laplacian 0.000008, dIdt 0.000884

10it [00:25, 2.53s/it]

Epoch 4, Epoch loss: total 462.159679, pixel 0.319729, grad 9.529771, laplacian 2812.653687, dIdt 17994.100098

Epoch 4, Epoch SSIM: pixel 0.056193, grad 0.020033, laplacian 0.000003, dIdt 0.001105

10it [00:25, 2.58s/it]

Epoch 5, Epoch loss: total 452.182953, pixel 0.319670, grad 8.941711, laplacian 2718.746960, dIdt 17941.377734

Epoch 5, Epoch SSIM: pixel 0.035697, grad 0.023118, laplacian 0.000003, dIdt 0.001593

10it [00:25, 2.54s/it]

Epoch 6, Epoch loss: total 353.849197, pixel 0.308936, grad 8.682603, laplacian 1736.287311, dIdt 17935.189844

Epoch 6, Epoch SSIM: pixel 0.077749, grad 0.021318, laplacian 0.000005, dIdt 0.000978

10it [00:25, 2.58s/it]

Epoch 7, Epoch loss: total 295.821411, pixel 0.312839, grad 8.332347, laplacian 1160.650034, dIdt 17892.285449

Epoch 7, Epoch SSIM: pixel 0.074073, grad 0.020851, laplacian 0.000004, dIdt 0.001449

10it [00:25, 2.51s/it]

Epoch 8, Epoch loss: total 368.288702, pixel 0.315093, grad 8.561647, laplacian 1888.020422, dIdt 17863.017383

Epoch 8, Epoch SSIM: pixel 0.058733, grad 0.018582, laplacian 0.000001, dIdt 0.000614

10it [00:26, 2.60s/it]

Epoch 9, Epoch loss: total 467.217441, pixel 0.319720, grad 8.641083, laplacian 2881.013983, dIdt 17825.162793

Epoch 9, Epoch SSIM: pixel 0.124318, grad 0.017580, laplacian 0.000002, dIdt 0.000401

10it [00:27, 2.80s/it]

Epoch 10, Epoch loss: total 713.758423, pixel 0.339931, grad 9.718257, laplacian 5347.353027, dIdt 17805.094336

Epoch 10, Epoch SSIM: pixel 0.186489, grad 0.015669, laplacian 0.000001, dIdt 0.000432

10it [00:25, 2.54s/it]

Epoch 11, Epoch loss: total 842.909290, pixel 0.325829, grad 10.048907, laplacian 6636.391064, dIdt 17826.495996

Epoch 11, Epoch SSIM: pixel 0.165996, grad 0.014063, laplacian 0.000001, dIdt 0.000339

10it [00:26, 2.62s/it]

Epoch 12, Epoch loss: total 1419.564450, pixel 0.368746, grad 11.414354, laplacian 12403.762891, dIdt 17804.634473

Epoch 12, Epoch SSIM: pixel 0.187385, grad 0.014755, laplacian 0.000001, dIdt 0.000213

10it [00:25, 2.52s/it]

Epoch 13, Epoch loss: total 1643.796716, pixel 0.365290, grad 12.368138, laplacian 14645.331445, dIdt 17802.634863

Epoch 13, Epoch SSIM: pixel 0.184158, grad 0.012082, laplacian 0.000001, dIdt 0.000158

10it [00:25, 2.57s/it]

Epoch 14, Epoch loss: total 1752.933618, pixel 0.381586, grad 12.499214, laplacian 15737.570801, dIdt 17792.616895

Epoch 14, Epoch SSIM: pixel 0.189889, grad 0.011957, laplacian 0.000000, dIdt 0.000182

10it [00:25, 2.51s/it]

Epoch 15, Epoch loss: total 1622.605798, pixel 0.404907, grad 11.703171, laplacian 14444.354102, dIdt 17699.966504

Epoch 15, Epoch SSIM: pixel 0.206009, grad 0.017348, laplacian 0.000000, dIdt 0.000115

10it [00:25, 2.56s/it]

Epoch 16, Epoch loss: total 1844.198645, pixel 0.483079, grad 12.193454, laplacian 16666.670264, dIdt 17631.177051

Epoch 16, Epoch SSIM: pixel 0.190710, grad 0.019222, laplacian 0.000000, dIdt 0.000139

10it [00:25, 2.52s/it]

Epoch 17, Epoch loss: total 2519.064728, pixel 0.548760, grad 14.309171, laplacian 23420.258789, dIdt 17560.734570

Epoch 17, Epoch SSIM: pixel 0.164187, grad 0.019163, laplacian 0.000000, dIdt 0.000040

10it [00:25, 2.51s/it]

Epoch 18, Epoch loss: total 3187.327197, pixel 0.652958, grad 17.297070, laplacian 30100.005762, dIdt 17559.616992

Epoch 18, Epoch SSIM: pixel 0.151996, grad 0.016163, laplacian 0.000000, dIdt 0.000067

10it [00:26, 2.61s/it]

Epoch 19, Epoch loss: total 2795.569495, pixel 0.672258, grad 17.069961, laplacian 26184.390332, dIdt 17542.270313



Epoch 19, Epoch SSIM: pixel 0.131105, grad 0.014221, laplacian 0.000000, dIdt 0.000086

10it [00:27, 2.79s/it]

Epoch 20, Epoch loss: total 3663.679553, pixel 0.769440, grad 18.601243, laplacian 34880.887598, dIdt 17372.982812

Epoch 20, Epoch SSIM: pixel 0.125092, grad 0.013949, laplacian 0.000000, dIdt 0.000175

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
/tmp/ipykernel_2988/1545087382.py in <module>
    55     for uniform_lr in learning_rates:
    56         model_path_full = model_path_act + 'uniformlr_' + "{:.0e}".
↳format(uniform_lr)
--> 57         run_siren(model_path_full, betas, total_epochs, [uniform_lr])
    58
    59     # 2. run with decaying learning rates

/tmp/ipykernel_2988/2675429778.py in run_siren(model_path, betas, total_epochs,
↳lr, cyclic, decay_exp, decay_multi)
     9         total_epochs=total_epochs, lr=lr,
    10         beta_0=betas[0], beta_1=betas[1], beta_2=betas[2],
↳beta_3=betas[3],
--> 11         cyclic=cyclic, decay_exp=decay_exp, decay_multi=decay_multi)
    12
    13     writer.close()

/tmp/ipykernel_2988/962146541.py in train(net, writer, img_path, niter,
↳total_epochs, lr, beta_0, beta_1, beta_2, beta_3, cyclic, decay_exp,
↳decay_multi)
    131     print("Epoch %d, Epoch SSIM: pixel %0.6f, grad %0.6f, laplacian
↳%0.6f, dIdt %0.6f" % (epoch, epoch_pixel_ssim/len(image), epoch_grad_ssim/
↳len(image), epoch_laplacian_ssim/len(image), epoch_dIdt_ssim/len(image)))
    132
--> 133     writer.add_graph(net, model_input)
    134     print("-----Finished-----")
    135

/opt/conda/lib/python3.7/site-packages/torch/utils/tensorboard/writer.py in
↳add_graph(self, model, input_to_model, verbose, use_strict_trace)
    734         if hasattr(model, 'forward'):
    735             # A valid PyTorch model should have a 'forward' method
--> 736             self._get_file_writer().add_graph(graph(model,
↳input_to_model, verbose, use_strict_trace))
    737         else:
    738             # Caffe2 models do not have the 'forward' method
```

```

/opt/conda/lib/python3.7/site-packages/torch/utils/tensorboard/_pytorch_graph.py
↳ in graph(model, args, verbose, use_strict_trace)
    289     with torch.onnx.select_model_mode_for_export(model, torch.onnx.
↳ TrainingMode.EVAL): # TODO: move outside of torch.onnx?
    290         try:
--> 291             trace = torch.jit.trace(model, args, strict=use_strict_trace)
    292             graph = trace.graph
    293             torch._C._jit_pass_inline(graph)

/opt/conda/lib/python3.7/site-packages/torch/jit/_trace.py in trace(func,
↳ example_inputs, optimize, check_trace, check_inputs, check_tolerance, strict,
↳ _force_outplace, _module_class, _compilation_unit)
    748         strict,
    749         _force_outplace,
--> 750         _module_class,
    751     )
    752

/opt/conda/lib/python3.7/site-packages/torch/jit/_trace.py in trace_module(mod,
↳ inputs, optimize, check_trace, check_inputs, check_tolerance, strict,
↳ _force_outplace, _module_class, _compilation_unit)
    989         _force_outplace,
    990         True,
--> 991         _module_class,
    992     )
    993     finally:

/opt/conda/lib/python3.7/site-packages/torch/autograd/grad_mode.py in
↳ decorate_context(*args, **kwargs)
    25     def decorate_context(*args, **kwargs):
    26         with self.clone():
---> 27             return func(*args, **kwargs)
    28     return cast(F, decorate_context)
    29

/opt/conda/lib/python3.7/site-packages/torch/jit/_trace.py in
↳ _check_trace(check_inputs, func, traced_func, check_tolerance, strict,
↳ _force_outplace, is_trace_module, _module_class)
    516     traced_outs = run_mod_and_filter_tensor_outputs(traced_func,
↳ inputs, "trace")
    517     fn_outs = run_mod_and_filter_tensor_outputs(func, inputs,
↳ "Python function")
--> 518     if compare_outputs(traced_outs, fn_outs, "Python function"):
    519         check_outs = run_mod_and_filter_tensor_outputs(
    520             check_mod_func, inputs, "repeated trace"

/opt/conda/lib/python3.7/site-packages/torch/jit/_trace.py in
↳ compare_outputs(original, reference, match_what)

```

```

495             rtol=check_tolerance,
496             atol=default_tolerances(orig, ref)[1],
--> 497             equal_nan=True,
498         )
499     except AssertionError as e:

[... skipping hidden 2 frame]

/opt/conda/lib/python3.7/site-packages/torch/testing/_comparison.py in
↳ compare(self)
    602         actual, expected = self._equalize_attributes(actual,
↳ expected)
    603
--> 604         self._compare_values(actual, expected)
    605
    606     @contextlib.contextmanager

/opt/conda/lib/python3.7/site-packages/torch/testing/_comparison.py in
↳ _compare_values(self, actual, expected)
    714         compare_fn = self._compare_regular_values_close
    715
--> 716         compare_fn(actual, expected, rtol=self.rtol, atol=self.atol,
↳ equal_nan=self.equal_nan)
    717
    718     def _compare_quantized_values(

/opt/conda/lib/python3.7/site-packages/torch/testing/_comparison.py in
↳ _compare_regular_values_close(self, actual, expected, rtol, atol, equal_nan,
↳ identifier)
    840         actual, expected = self._promote_for_comparison(actual, expected)
    841         matches = torch.isclose(actual, expected, rtol=rtol, atol=atol,
↳ equal_nan=equal_nan)
--> 842         if torch.all(matches):
    843             return
    844

KeyboardInterrupt:

```

Error in callback <function install\_repl\_displayhook.<locals>.post\_execute at 0x7fd3ac486ef0> (for post\_execute):

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
/opt/conda/lib/python3.7/site-packages/matplotlib/pyplot.py in post_execute()
    136     def post_execute():
    137         if matplotlib.is_interactive():
--> 138             draw_all()

```

```

139
140     try: # IPython >= 2

/opt/conda/lib/python3.7/site-packages/matplotlib/_pylab_helpers.py in
↳draw_all(cls, force)
    135         for manager in cls.get_all_fig_managers():
    136             if force or manager.canvas.figure.stale:
--> 137                 manager.canvas.draw_idle()
    138
    139

/opt/conda/lib/python3.7/site-packages/matplotlib/backend_bases.py in
↳draw_idle(self, *args, **kwargs)
    2058         if not self._is_idle_drawing:
    2059             with self._idle_draw_cntx():
-> 2060                 self.draw(*args, **kwargs)
    2061
    2062     @property

/opt/conda/lib/python3.7/site-packages/matplotlib/backends/backend_agg.py in
↳draw(self)
    434         (self.toolbar._wait_cursor_for_draw_cm() if self.toolbar
    435          else nullcontext()):
--> 436         self.figure.draw(self.renderer)
    437         # A GUI class may be need to update a window using this
↳draw, so
    438         # don't forget to call the superclass.

/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in
↳draw_wrapper(artist, renderer, *args, **kwargs)
    71     @wraps(draw)
    72     def draw_wrapper(artist, renderer, *args, **kwargs):
----> 73         result = draw(artist, renderer, *args, **kwargs)
    74         if renderer._rasterizing:
    75             renderer.stop_rasterizing()

/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in
↳draw_wrapper(artist, renderer)
    48         renderer.start_filter()
    49
----> 50         return draw(artist, renderer)
    51     finally:
    52         if artist.get_agg_filter() is not None:

/opt/conda/lib/python3.7/site-packages/matplotlib/figure.py in draw(self,
↳renderer)
    2809         self.patch.draw(renderer)
    2810         mimage._draw_list_compositing_images(

```

```

-> 2811             renderer, self, artists, self.suppressComposite)
    2812
    2813             for sfig in self.subfigs:

/opt/conda/lib/python3.7/site-packages/matplotlib/image.py in
-> _draw_list_compositing_images(renderer, parent, artists, suppress_composite)
    130         if not_composite or not has_images:
    131             for a in artists:
--> 132                 a.draw(renderer)
    133         else:
    134             # Composite any adjacent images together

/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in
-> draw_wrapper(artist, renderer)
    48             renderer.start_filter()
    49
---> 50             return draw(artist, renderer)
    51         finally:
    52             if artist.get_agg_filter() is not None:

/opt/conda/lib/python3.7/site-packages/matplotlib/axes/_base.py in draw(self,
-> renderer)
    3081
    3082         mimage._draw_list_compositing_images(
-> 3083             renderer, self, artists, self.figure.suppressComposite)
    3084
    3085         renderer.close_group('axes')

/opt/conda/lib/python3.7/site-packages/matplotlib/image.py in
-> _draw_list_compositing_images(renderer, parent, artists, suppress_composite)
    130         if not_composite or not has_images:
    131             for a in artists:
--> 132                 a.draw(renderer)
    133         else:
    134             # Composite any adjacent images together

/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in
-> draw_wrapper(artist, renderer)
    48             renderer.start_filter()
    49
---> 50             return draw(artist, renderer)
    51         finally:
    52             if artist.get_agg_filter() is not None:

/opt/conda/lib/python3.7/site-packages/matplotlib/axis.py in draw(self,
-> renderer, *args, **kwargs)
    1161

```

```

1162         for tick in ticks_to_draw:
-> 1163             tick.draw(renderer)
1164
1165             # scale up the axis label box to also find the neighbors, not

/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in
->draw_wrapper(artist, renderer)
    48         renderer.start_filter()
    49
---> 50         return draw(artist, renderer)
    51     finally:
    52         if artist.get_agg_filter() is not None:

/opt/conda/lib/python3.7/site-packages/matplotlib/axis.py in draw(self, rendere.)
    297         for artist in [self.gridline, self.tick1line, self.tick2line,
    298                        self.label1, self.label2]:
-> 299             artist.draw(renderer)
    300         renderer.close_group(self.__name__)
    301         self.stale = False

/opt/conda/lib/python3.7/site-packages/matplotlib/artist.py in
->draw_wrapper(artist, renderer)
    48         renderer.start_filter()
    49
---> 50         return draw(artist, renderer)
    51     finally:
    52         if artist.get_agg_filter() is not None:

/opt/conda/lib/python3.7/site-packages/matplotlib/text.py in draw(self, rendere.)
    682         renderer.open_group('text', self.get_gid())
    683
-> 684         with self._cm_set(text=self._get_wrapped_text()):
    685             bbox, info, descent = self._get_layout(renderer)
    686             trans = self.get_transform()

/opt/conda/lib/python3.7/contextlib.py in __enter__(self)
    105         return self.__class__(self.func, self.args, self.kwds)
    106
-> 107     def __enter__(self):
    108         # do not keep args and kwds alive unnecessarily
    109         # they are only needed for recreation, which is not possible
-> anymore

KeyboardInterrupt:

```

Error in callback <function flush\_figures at 0x7fd3acbd4830> (for post\_execute):

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
/opt/conda/lib/python3.7/site-packages/matplotlib_inline/backend_inline.py in
↳ flush_figures()
    119         # ignore the tracking, just draw and close all figures
    120         try:
--> 121             return show(True)
    122         except Exception as e:
    123             # safely show traceback if in IPython, else raise

/opt/conda/lib/python3.7/site-packages/matplotlib_inline/backend_inline.py in
↳ show(close, block)
    41         display(
    42             figure_manager.canvas.figure,
---> 43             metadata=_fetch_figure_metadata(figure_manager.canvas.
↳ figure)
    44         )
    45         finally:

/opt/conda/lib/python3.7/site-packages/IPython/core/display.py in
↳ display(include, exclude, metadata, transient, display_id, *objs, **kwargs)
    318         publish_display_data(data=obj, metadata=metadata, **kwargs)
    319         else:
--> 320         format_dict, md_dict = format(obj, include=include,
↳ exclude=exclude)
    321         if not format_dict:
    322             # nothing to display (e.g. _ipython_display_ took over)

/opt/conda/lib/python3.7/site-packages/IPython/core/formatters.py in
↳ format(self, obj, include, exclude)
    178         md = None
    179         try:
--> 180         data = formatter(obj)
    181         except:
    182             # FIXME: log the exception

/opt/conda/lib/python3.7/site-packages/decorator.py in fun(*args, **kw)
    230         if not kwsyntax:
    231             args, kw = fix(args, kw, sig)
--> 232         return caller(func, *(extras + args), **kw)
    233         fun.__name__ = func.__name__
    234         fun.__doc__ = func.__doc__

/opt/conda/lib/python3.7/site-packages/IPython/core/formatters.py in
↳ catch_format_error(method, self, *args, **kwargs)
    222         """show traceback on failed format call"""
    223         try:

```

```

--> 224         r = method(self, *args, **kwargs)
      225     except NotImplementedError:
      226         # don't warn on NotImplementedError

/opt/conda/lib/python3.7/site-packages/IPython/core/formatters.py in
-> __call__(self, obj)
      339         pass
      340     else:
--> 341         return printer(obj)
      342         # Finally look for special method names
      343     method = get_real_method(obj, self.print_method)

/opt/conda/lib/python3.7/site-packages/IPython/core/pylabtools.py in
-> print_figure(fig, fmt, bbox_inches, base64, **kwargs)
      149     FigureCanvasBase(fig)
      150
--> 151     fig.canvas.print_figure(bytes_io, **kw)
      152     data = bytes_io.getvalue()
      153     if fmt == 'svg':

/opt/conda/lib/python3.7/site-packages/matplotlib/backend_bases.py in
-> print_figure(self, filename, dpi, facecolor, edgecolor, orientation, format,
-> bbox_inches, pad_inches, bbox_extra_artists, backend, **kwargs)
      2298         if bbox_inches == "tight":
      2299             bbox_inches = self.figure.get_tightbbox(
-> 2300                 renderer, bbox_extra_artists=bbox_extra_artists

      2301             if pad_inches is None:
      2302                 pad_inches = rcParams['savefig.pad_inches']

/opt/conda/lib/python3.7/site-packages/matplotlib/figure.py in
-> get_tightbbox(self, renderer, bbox_extra_artists)
      1630
      1631     for a in artists:
-> 1632         bbox = a.get_tightbbox(renderer)
      1633         if bbox is not None and (bbox.width != 0 or bbox.height !=
-> 0):
      1634             bb.append(bbox)

/opt/conda/lib/python3.7/site-packages/matplotlib/axes/_base.py in
-> get_tightbbox(self, renderer, call_axes_locator, bbox_extra_artists,
-> for_layout_only)
      4627         try:
      4628             bb_yaxis = self.yaxis.get_tightbbox(
-> 4629                 renderer, for_layout_only=for_layout_only)

      4630         except TypeError:
      4631             # in case downstream library has redefined axis:

```



```

/opt/conda/lib/python3.7/site-packages/matplotlib/axis.py in get_tightbbox(self,
↳renderer, for_layout_only)
    1103         ticks_to_draw = self._update_ticks()
    1104
-> 1105         self._update_label_position(renderer)
    1106
    1107         # go back to just this axis's tick labels

/opt/conda/lib/python3.7/site-packages/matplotlib/axis.py in
↳_update_label_position(self, renderer)
    2350         # get bounding boxes for this axis and any siblings
    2351         # that have been set by `fig.align_ylabels()`
-> 2352         bboxes, bboxes2 = self.
↳_get_tick_boxes_siblings(renderer=renderer)
    2353
    2354         x, y = self.label.get_position()

/opt/conda/lib/python3.7/site-packages/matplotlib/axis.py in
↳_get_tick_boxes_siblings(self, renderer)
    1878         for ax in grouper.get_siblings(self.axes):
    1879             axis = getattr(ax, f"{axis_name}axis")
-> 1880             ticks_to_draw = axis._update_ticks()
    1881             tlb, tlb2 = axis._get_tick_bboxes(ticks_to_draw, renderer)
    1882             bboxes.extend(tlb)

/opt/conda/lib/python3.7/site-packages/matplotlib/axis.py in _update_ticks(self
    1046         major_labels = self.major.formatter.format_ticks(major_locs)
    1047         major_ticks = self.get_major_ticks(len(major_locs))
-> 1048         self.major.formatter.set_locs(major_locs)
    1049         for tick, loc, label in zip(major_ticks, major_locs,
↳major_labels):
    1050             tick.update_position(loc)

/opt/conda/lib/python3.7/site-packages/matplotlib/ticker.py in set_locs(self,
↳locs)
    709         if self._useOffset:
    710             self._compute_offset()
--> 711         self._set_order_of_magnitude()
    712         self._set_format()
    713

/opt/conda/lib/python3.7/site-packages/matplotlib/ticker.py in
↳_set_order_of_magnitude(self)
    765         vmin, vmax = sorted(self.axis.get_view_interval())
    766         locs = np.asarray(self.locs)
--> 767         locs = locs[(vmin <= locs) & (locs <= vmax)]
    768         locs = np.abs(locs)
    769         if not len(locs):

```

KeyboardInterrupt:

### 0.0.12 Visualize Results

```
[28]: %load_ext tensorboard
      %tensorboard --logdir="runs"
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6006 (pid 2841), started 0:11:02 ago. (Use '!kill\_↵  
↵2841' to kill it.)

<IPython.core.display.HTML object>

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```